



# ECE408 /CS483/CSE408 Fall 2024

## Applied Parallel Programming

### Lecture 2:

# Introduction to CUDA C and Data Parallel Programming

# Course Reminders

- Lab 0 is due on Friday September 6th at 8pm US Central time
  - Should be released soon, we will notify you when this happens.
  - It is an easy lab, but it may take some time to get the tools in place.
  - Its main purpose is to get familiar with the programming environment and the process.
  - If you miss this deadline, that's OK for Lab 0, but you must submit it anyway. Remember, Lab 0 will be graded, but not counted towards your overall grade.

# Objective

- To learn the basic concept of data parallel computing
- To learn the basic features of the CUDA C programming interface


# Thread as a basic unit of computing

- What is a thread?


- Program
- PC
- Context
  - Memory
  - Registers
  - ...




- Multiple threads




- Many threads




# *A Data Parallel Computation Example: Conversion of a color image to grey-scale image*



```
for each pixel {  
    pixel = gsConvert(pixel)  
}  
// Every pixel is independent  
// of every other pixel
```




*The pixels can be calculated independently of each other*



```
for each pixel {  
    pixel = gsConvert(pixel)  
}  
// Every pixel is independent  
// of every other pixel
```


# CUDA/OpenCL – Execution Model

- Integrated host+device app C program
  - Serial or modestly parallel parts in **host** C code
  - Highly parallel parts in **device** SPMD kernel C code




# Arrays of Parallel Threads

- A CUDA kernel is executed as a **grid** (array) of threads
  - All threads in a grid run the same kernel code
  - Single Program Multiple Data (SPMD model)
  - Each thread has **a unique index** that it uses to compute memory addresses and make control decisions




# Logical Execution Model for CUDA

- Each CUDA kernel
  - is executed by a **grid**,
  - a 3D array of **thread blocks**, which are
  - 3D arrays of **threads**.




# Single Program, Multiple Data

- Each thread
  - executes the **same program**
  - on **distinct data inputs**,
  - a single-program, multiple-data (**SPMD**) model



# gridDim Gives Number of Blocks


- Number of blocks in each dimension is
  - `gridDim.x` ... 8
  - `gridDim.y` ... 3
  - `gridDim.z` ... 2



For 2D (and 1D grids), simply use grid dimension 1 for Z (and Y).


# blockIdx is Unique for Each Block

- Each block has a unique index tuple
  - `blockIdx.x` (from 0 to  $(\text{gridDim.x} - 1)$  )
  - `blockIdx.y` (from 0 to  $(\text{gridDim.y} - 1)$  )
  - `blockIdx.z` (from 0 to  $(\text{gridDim.z} - 1)$  )



# blockDim: # of Threads per Block


- Number of blocks in each dimension is
  - `blockDim.x` ... **5**
  - `blockDim.y` ... **4**
  - `blockDim.z` ... **3**



For 2D (and 1D blocks), simply use block dimension 1 for Z (and Y).

# threadIdx Unique for Each Thread


- Each thread has a unique index tuple
  - `threadIdx.x` (from 0 to  $(\text{blockDim.x} - 1)$  )
  - `threadIdx.y` (from 0 to  $(\text{blockDim.y} - 1)$  )
  - `threadIdx.z` (from 0 to  $(\text{blockDim.z} - 1)$  )



threadIdx tuple is unique to each thread  
**WITHIN A BLOCK.**


# Thread Blocks: Scalable Cooperation

- Threads within a block cooperate via **shared memory, atomic operations** and **barrier synchronization** (to be covered later)
- Threads in different blocks cooperate less.




# blockIdx and threadIdx

- Thread block and thread organization
  - simplifies memory addressing
  - when processing multidimensional data
- Image processing
- Vectors, matrices, tensors
- Solving PDEs on volumes
- ...



# Vector Addition – Conceptual View



# Vector Addition – Traditional C Code

```
// Compute vector sum C = A+B
void vecAdd(float* A, float* B, float* C, int n)
{
    for (i = 0, i < n, i++)
        C[i] = A[i] + B[i];
}

int main()
{
    // Memory allocation for A_h, B_h, and C_h
    // I/O to read A_h and B_h, N elements
    ...
    vecAdd(A_h, B_h, C_h, N);
}
```

# Heterogeneous Computing: vecAdd Host Code

```
#include <cuda.h>
void vecAdd(float* A, float* B, float* C, int n)
{
    int size = n* sizeof(float);
    float *A_d, *B_d, *C_d;
    ...
1. // Allocate device memory for A, B, and C
   // copy A and B to device memory


2. // Kernel launch code - to have the device
   // to perform the actual vector addition

3. // copy C from the device memory
   // Free device vectors
}
```

# Partial Overview of CUDA Memories


- Device code can:
  - R/W per-thread **registers**
  - R/W per-grid **global memory**
- Host code can
  - Transfer data to/from per grid **global memory**

We will cover more later.




# CUDA Device Memory Management API functions

- `cudaMalloc()`
  - Allocates object in the device global memory
  - Two parameters
    - **Address of a pointer** to the allocated object
    - **Size of** the allocated object in terms of bytes
- `cudaFree()`
  - Frees object from device global memory
    - **Pointer** to freed object



# Host-Device Data Transfer API functions

- `cudaMemcpy()`
  - memory data transfer
  - Requires four parameters
    - Pointer to destination
    - Pointer to source
    - Number of bytes copied
    - Type/Direction of transfer



```
void vecAdd(float* A, float* B, float* C, int n)
{
    int size = n * sizeof(float);
    float *A_d, *B_d, *C_d;

1. // Transfer A and B to device memory
    // (error-checking omitted)
    cudaMalloc((void **) &A_d, size);
cudaMemcpy(A_d, A, size, cudaMemcpyHostToDevice);
    cudaMalloc((void **) &B_d, size);
cudaMemcpy(B_d, B, size, cudaMemcpyHostToDevice);

    // Allocate device memory for
    cudaMalloc((void **) &C_d, size);

2. // Kernel invocation code - to be shown later
    ...
3. // Transfer C from device to host
cudaMemcpy(C, C_d, size, cudaMemcpyDeviceToHost);
    // Free device memory for A, B, C
    cudaFree(A_d); cudaFree(B_d); cudaFree(C_d);
}
```

# Example: Vector Addition Kernel

Device Code

```
// Compute vector sum C = A+B
// Each thread performs one pair-wise addition
__global
void vecAddKernel(float* A_d, float* B_d, float* C_d, int n)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x ;
    if(i<n) C_d[i] = A_d[i] + B_d[i];
}

int vectAdd(float* A, float* B, float* C, int n)
{
    // A_d, B_d, C_d allocations and copies omitted
    // Run ceil(n/256) blocks of 256 threads each
    vecAddKernel<<<ceil(n/256.0), 256>>>(A_d, B_d, C_d, n);
}
```

# Example: Vector Addition Kernel

```
// Compute vector sum C = A+B
// Each thread performs one pair-wise addition
__global__
void vecAddKernel(float* A_d, float* B_d, float* C_d, int n)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if(i<n) C_d[i] = A_d[i] + B_d[i];
}

int vecAdd(float* A, float* B, float* C, int n)
{
    // A_d, B_d, C_d allocations and copies omitted
    // Run ceil(n/256) blocks of 256 threads each
    vecAddKernel<<<ceil(n/256.0),256>>>(A_d, B_d, C_d, n);
}
```

Host Code

# More on Kernel Launch

## Equivalent Host Code

```
int vecAdd(float* A, float* B, float* C, int n)
{
    // A_d, B_d, C_d allocations and copies omitted
    // Run ceil(n/256) blocks of 256 threads each
    dim3 DimGrid(n/256, 1, 1);
    if (0 != (n % 256)) { DimGrid.x++; }
    dim3 DimBlock(256, 1, 1);

    vecAddKernel<<<DimGrid,DimBlock>>>(A_d, B_d, C_d, n);
}
```

- Any call to a kernel function is asynchronous from CUDA 1.0 on, explicit synch needed for blocking

# Vector Addition Kernel

```
// Compute vector sum C = A+B
// Each thread performs one pair-wise addition
__global__
void vecAddKernel(float* A_d, float* B_d, float* C_d, int n)
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    if(i<n) C_d[i] = A_d[i] + B_d[i];
}

int vecAdd(float* A, float* B, float* C, int n)
{
    // A_d, B_d, C_d allocations and copies omitted
    // Run ceil(n/256) blocks of 256 threads each
    dim3 DimGrid(ceil(n/256), 1, 1);
    dim3 DimBlock(256, 1, 1);
    vecAddKernel<<<DimGrid,DimBlock>>>(A_d, B_d, C_d, n);
}
```

**A** Number of blocks per dimension

**B** Number of threads per dimension in a block

**C** Unique block # in x dimension

**D** Number of threads per block in x dimension

**E** Unique thread # in x dimension in the block


# Kernel execution in a nutshell

```
__host__
void vecAdd()
{
    dim3 DimGrid(ceil(n/256.0),1,1);
    dim3 DimBlock(256,1,1);

    vecAddKernel<<<DimGrid,DimBlock>>>
    (A_d,B_d,C_d,n);
}
```

```
__global__
void vecAddKernel(float *A_d,
float *B_d, float *C_d, int n)
{
    int i = blockIdx.x * blockDim.x
            + threadIdx.x;

    if( i < n ) C_d[i] = A_d[i]+B_d[i];
}
```




# More on CUDA Function Declarations

	Executed on the:	Only callable from the:
<code>__device__ float DeviceFunc()</code>	device	device
<code>__global__ void KernelFunc()</code>	device	host
<code>__host__ float HostFunc()</code>	host	host

- `__global__` defines a kernel function
  - Each “`__`” consists of two underscore characters
  - A kernel function must return `void`
- `__device__` and `__host__` can be used together

# Compiling A CUDA Program

Integrated C programs with CUDA extensions





**ANY MORE QUESTIONS?  
READ CHAPTER 2**

# Problem Solving

- Consider the following code:

```
kernel<<VECTOR_N, ELEMENT_N>>>(d_C, d_A, d_B, ELEMENT_N);
```

- Q: How many CUDA threads are in each block as the result of the following kernel call?
- A: **ELEMENT\_N**
- Q: How many CUDA threads will be created as the result of the following kernel call?
- A: **VECTOR\_N \* ELEMENT\_N**

# Problem Solving

- Q: For a vector addition, assume that the vector length is 16000, each thread calculates 8 output elements, and the thread block size is 256 threads. The programmer configures the kernel launch to have a minimal number of thread blocks to cover all output elements. How many **threads** will be in the **grid**?
- A:
  - How many threads do we need?  $16000/8 = 2000$
  - How many blocks of threads do we need to run 2000 threads?  
 $\text{ceil}(2000/256) = 8$
  - Thus, how many threads will be running?  $8 * 256 = 2048$

# Problem Solving

- Q: A CUDA kernel is launched with 512 thread blocks each of which has 256 threads. If a variable is declared as a local variable in the kernel, how many versions of the variable will be created through the lifetime of the execution of the kernel?
- A:
  - How many threads will be created?  $512 * 256 = 131072$
  - So, there will be as many copies of the local variable, one in each thread.