



ECE408/CS483/CSE408 Fall 2024

# Applied Parallel Programming

## Lecture 9: Tiled Convolution Analysis

# Course Reminders

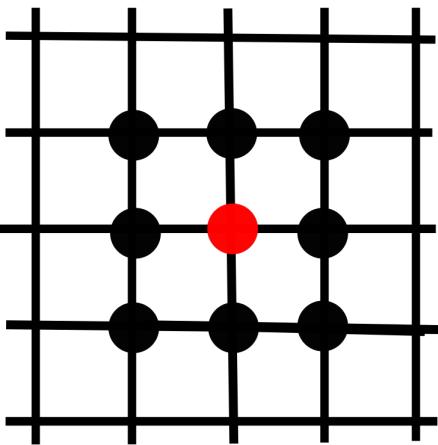
- Lab 1 is graded, 2 will be graded very soon
  - Check \_grade branch of your github repo for feedback
  - There is a regrade request procedure, see posts from the Lab TA
- Lab 3 is due this Friday
- Midterm 1 is on Tuesday, October 15<sup>th</sup>
  - See Canvas for details
- Project Milestone 1: Baseline CPU & GPU implementation
  - Project details to be posted soon

# Objective

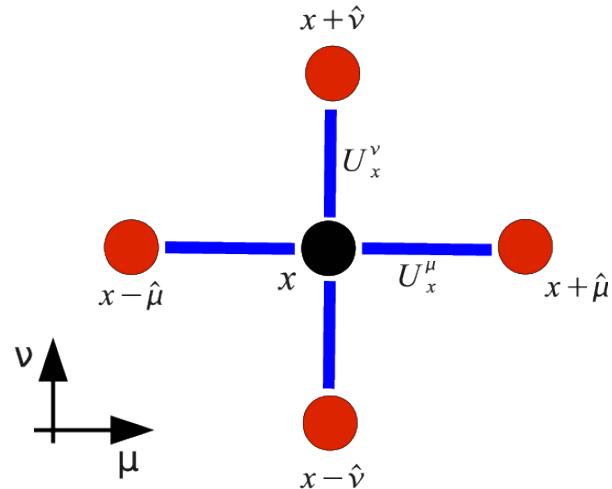
- To learn more about the analysis of tiled convolution/stencil algorithms

# Stencil Algorithms

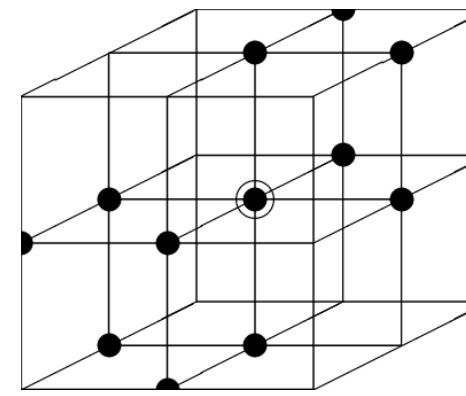
- Numerical data processing algorithms which update array elements according to some fixed pattern, called a *stencil*
  - Convolution is just one such example



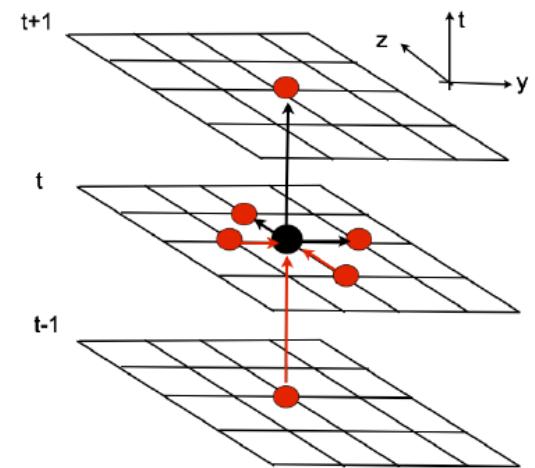
2D convolutional  
kernel (9-point 2D  
stencil)



Nearest neighbor lattice  
Dirac operator (5-point 2D  
stencil)

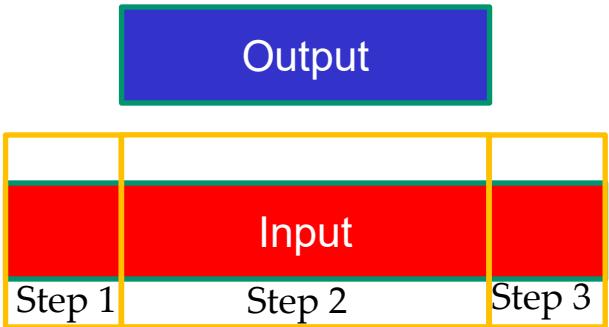


Finite difference stencil  
for 3D explicit time-  
marching (13-point 3D  
stencil)



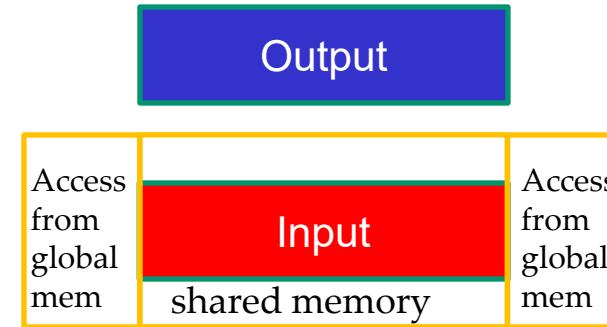
The Wilson-Dslash  
operator  
(4D stencil)

# Review: Three Tiling Strategies



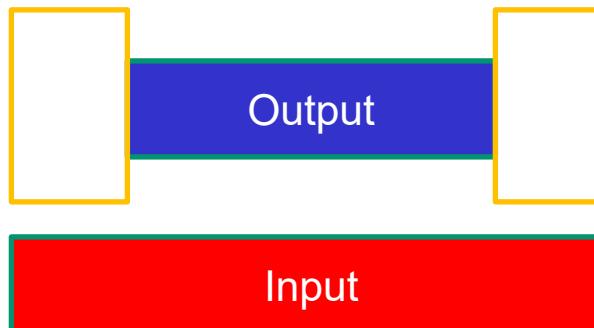
Strategy 1

1. Block size covers **output** tile
2. Use multiple steps to load input tile



Strategy 3

1. Block size covers **output** tile
2. Load only “core” of input tile
3. Access halo cells from global memory

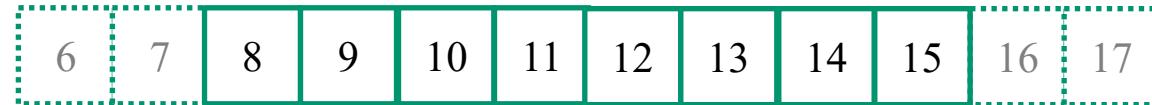


Strategy 2

1. Block size covers **input** tile
2. Load input tile in one step
3. Turn off some threads when calculating output

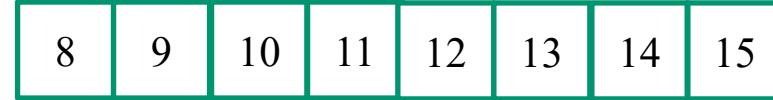
# A Small 1D Convolution Example

tile



MASK\_WIDTH is 5

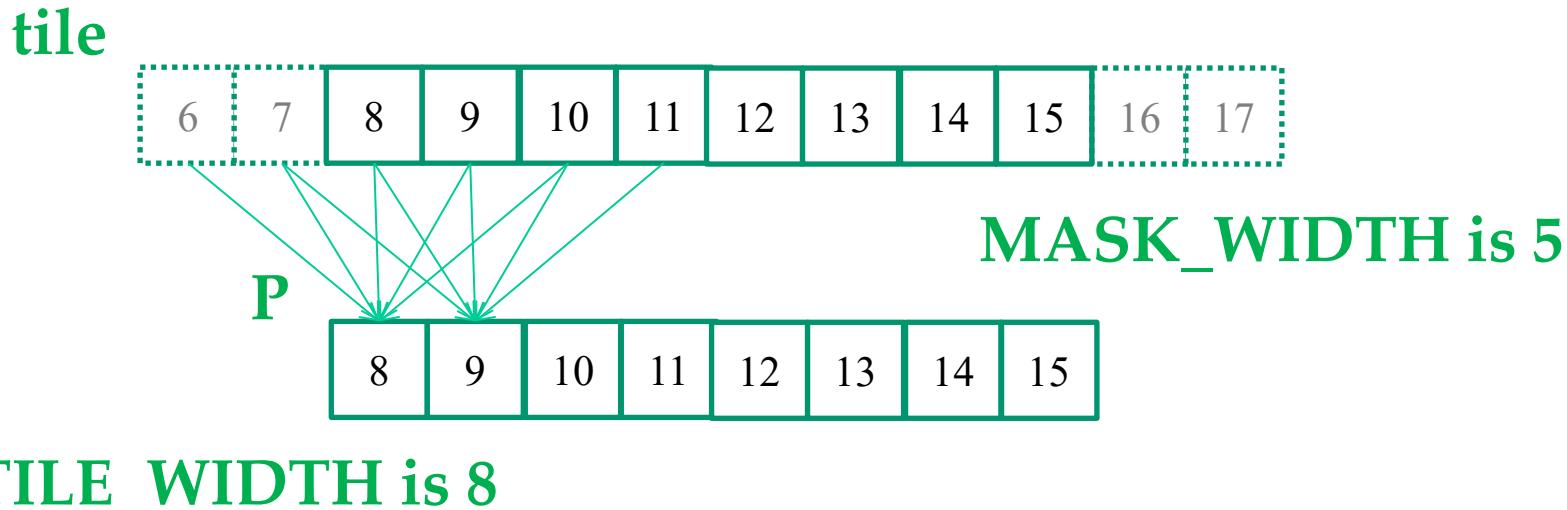
P



TILE\_WIDTH is 8

- output and input tiles for block 1
- For MASK\_WIDTH of 5, each block loads  $8 + (5 - 1) = 12$  elements (**12 memory loads**)

# Each Output Uses MASK\_WIDTH Inputs



- $P[8]$  uses  $N[6], N[7], N[8], N[9], N[10]$
- $P[9]$  uses  $N[7], N[8], N[9], N[10], N[11]$
- ...
- $P[15]$  uses  $N[13], N[14], N[15], N[16], N[17]$

Total of **8 \* 5 values** from tile used for the output.

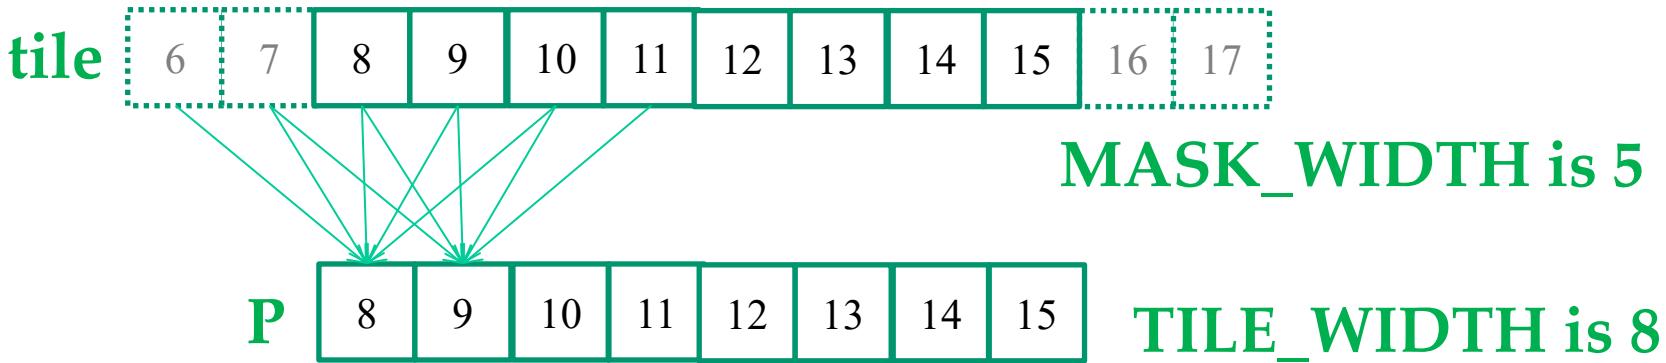
# A simple way to calculate tiling benefit

- $8+(5-1) = 12$  unique elements of input array N loaded
- $8*5$  global memory accesses potentially replaced by shared memory accesses
- This gives a bandwidth reduction of  **$40/12=3.3$**
- This is independent of the size of N

# In General, for 1D convolution kernels

- Load  $(\text{TILE\_WIDTH} + \text{MASK\_WIDTH} - 1)$  elements from global memory to shared memory
- Replace  $(\text{TILE\_WIDTH} * \text{MASK\_WIDTH})$  global memory accesses with shared memory accesses
- This leads to bandwidth reduction of  
 **$(\text{TILE\_SIZE} * \text{MASK\_WIDTH}) / (\text{TILE\_SIZE} + \text{MASK\_WIDTH} - 1)$**

# Another Way to Look at Reuse



- tile[6] is used by P[8] (1×)
- tile[7] is used by P[8], P[9] (2×)
- tile[8] is used by P[8], P[9], P[10] (3×)
- tile[9] is used by P[8], P[9], P[10], P[11] (4×)
- tile[10] is used by P[8], P[9], P[10], P[11], P[12] (5×)
- ... (5×)
- tile[14] is used by P[12], P[13], P[14], P[15] (4×)
- tile[15] is used by P[13], P[14], P[15] (3×)
- tile[16] is used by P[14], P[15] (2×)
- tile[17] is used by P[15] (1×)

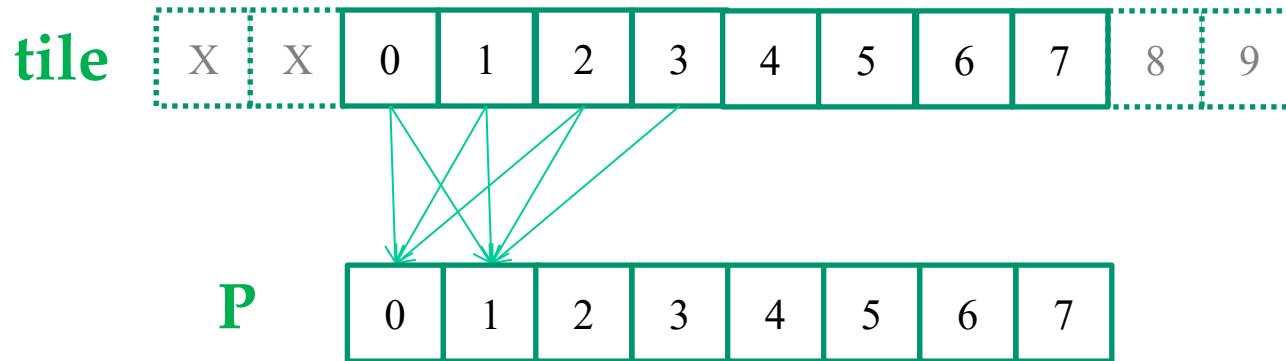
# Another Way to Look at Reuse

- Each access **tile** replaces an access to input N
- The total number of global memory accesses (to the  $(8+5-1)=12$  input elements) replaced by shared memory accesses is

$$\begin{aligned} & 1 + 2 + 3 + 4 + 5 * (8-5+1) + 4 + 3 + 2 + 1 \\ & = 10 + 20 + 10 \\ & = 40 \end{aligned}$$

- There are 12 elements of input N, so the average reduction is  
**40/12 = 3.3**

# What about Boundary Tiles?



- $P[0]$  uses  $N[0], N[1], N[2]$
- $P[1]$  uses  $N[0], N[1], N[2], N[3]$
- $P[2]$  uses  $N[0], N[1], N[2], N[3], N[4]$

Less than  $8 * 5$  elements of  $N$  are used for the output tile

# Ghost elements change ratios

- For a boundary tile, we load  $\text{TILE\_WIDTH} + (\text{MASK\_WIDTH}-1)/2$  elements
  - 10 in our example of  $\text{TILE\_WIDTH}$  of 8 and  $\text{MASK\_WIDTH}$  of 5
- Computing boundary elements do not access global memory for ghost cells
  - Total accesses is  $6*5 + 4 + 3 = 37$  accesses (when computing the P elements)

The reduction is **37/10 = 3.7**

# In General, for 1D, internal tiles

- The total number of global memory accesses to the  $(TILE\_WIDTH+Mask\_Width-1)$  elements of input N replaced by shared memory accesses is

$$\begin{aligned} & 1 + 2 + \dots + Mask\_Width-1 + Mask\_Width * (TILE\_WIDTH - Mask\_Width + 1) + Mask\_Width-1 + \dots + 2 + 1 \\ &= ((Mask\_Width-1) * Mask\_Width)/2 + Mask\_Width * (TILE\_WIDTH - Mask\_Width + 1) + ((Mask\_Width-1) * Mask\_Width)/2 \\ &= (Mask\_Width-1) * Mask\_Width + Mask\_Width * (TILE\_WIDTH - Mask\_Width + 1) \\ &= Mask\_Width * TILE\_WIDTH \end{aligned}$$

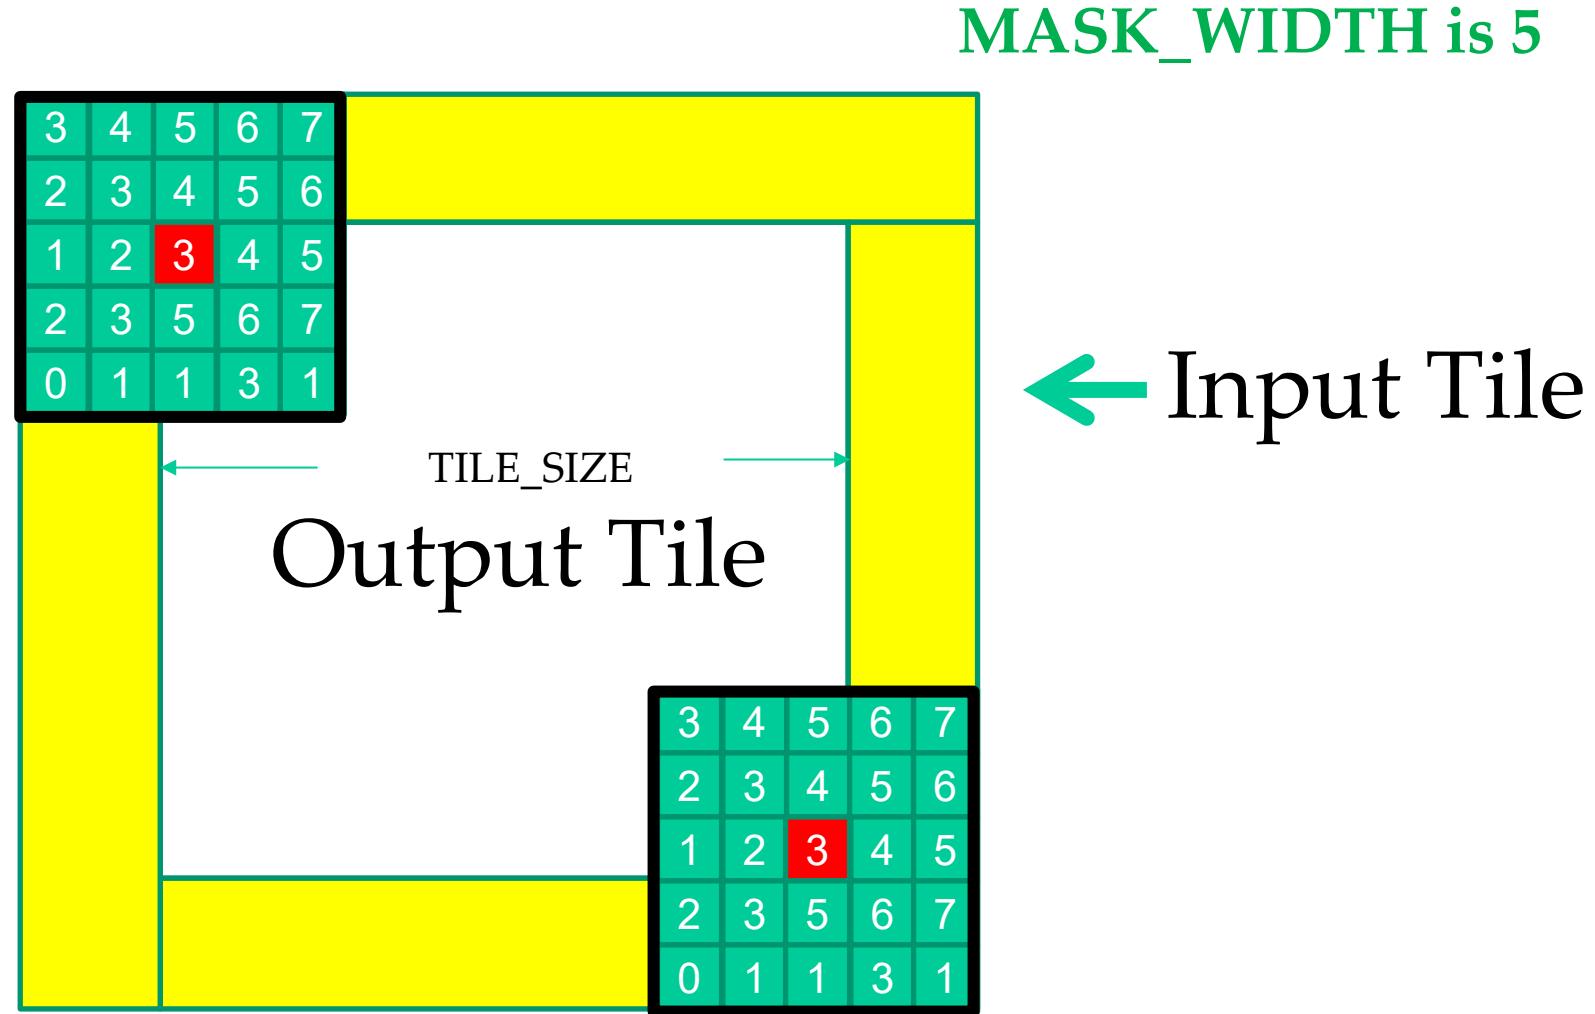
# Bandwidth Reduction for 1D

- The reduction is

$$(\text{TILE\_SIZE} * \text{MASK\_WIDTH}) / (\text{TILE\_SIZE} + \text{MASK\_WIDTH} - 1)$$

TILE_WIDTH	16	32	64	128	256
Reduction MASK_WIDTH = 5	4.0	4.4	4.7	4.9	4.9
Reduction MASK_WIDTH = 9	6.0	7.2	8.0	8.5	8.7

# Review: Parallelization of Tile Load



# Analysis for an 8×8 Output Tile, MASK\_WIDTH of 5

- Loading input tile requires **(8+5-1)<sup>2</sup> = 144 reads**
- Calculation of each output requires  $5^2 = 25$  input elements
- **8×8×25 = 1,600 global memory accesses** for computing output tile are converted to shared memory accesses
- Bandwidth reduction of **1,600/144 = 11.1×**

# In General

- $(\text{TILE\_WIDTH} + \text{MASK\_WIDTH} - 1)^2$  elements need to be loaded from  $\mathbf{N}$  into shared memory
- The calculation of each  $\mathbf{P}$  element needs to access  $\text{MASK\_WIDTH}^2$  elements of  $\mathbf{N}$
- $(\text{TILE\_WIDTH} * \text{MASK\_WIDTH})^2$  global memory accesses converted into shared memory accesses
- Bandwidth reduction of  
$$(\text{TILE\_WIDTH} * \text{MASK\_WIDTH})^2 / (\text{TILE\_WIDTH} + \text{MASK\_WIDTH} - 1)^2$$

# Bandwidth Reduction for 2D Convolution Kernel

- The reduction is

$$(\text{ TILE\_WIDTH} * \text{ MASK\_WIDTH})^2 / (\text{ TILE\_WIDTH} + \text{ MASK\_WIDTH} - 1)^2$$

TILE_WIDTH	8	16	32	64
Reduction MASK_WIDTH = 5	11.1	16	19.7	22.1
Reduction MASK_WIDTH= 9	20.3	36	51.8	64

# Ghost elements change ratios

- 2D calculation left for your enjoyment

# 2B/FLOP for Untiled Convolution

- How much global memory per FLOP is in untiled convolution?
- In untiled convolution,
  - each value from **N** (**4B** from global memory)
  - is multiplied by a value from **M** (**4B** from constant cache, **1 FLOP**),
  - then added to a running sum (**1 FLOP**)
- That gives **2B / FLOP**

# Full Use of Compute Requires 13.3× Reuse

- Recall our reuse discussion from matrix multiply:
  - 1,000 **GFLOP/s** for GPU from ~2010, and
  - 150 **GB/s** memory bandwidth.
- Dividing memory bandwidth by **2B/FLOP**,
$$\frac{150 \text{ GB/s}}{2 \text{ B/FLOP}} = 75 \text{ GFLOP/s} = 7.50\% \text{ of peak.}$$
- **Need at least  $100/7.50 = 13.3\times$  reuse** to make full use of compute resources

# In 2020, Need 52.1× Reuse

- In 2020, the **GRID K520** offers
  - nearly **5,000 GFLOP/s**, but only
  - **192 GB/s** memory bandwidth
- Dividing memory bandwidth by **2B/FLOP**,  
$$\frac{192 \text{ GB/s}}{2 \text{ B/FLOP}} = \mathbf{96 \text{ GFLOP/s} = 1.92\% \text{ of peak}}$$
- **Need at least  $100/1.92 = 52.1\times$  reuse** to make full use of compute resources

# Need ~26× Reuse on H100 GPUs

- In 2023, the H100 PCIe version offers
  - 26 TFLOP/s, and
  - 2 TB/s memory bandwidth
- Dividing memory bandwidth by 2B/FLOP,  
$$\frac{2 \text{ TB/s}}{2 \text{ B/FLOP}} = 1 \text{ TFLOP/s} = 3.85\% \text{ of peak}$$
- Need at least  $100/3.85 = 25.97\times$  reuse to make full use of compute resources

# Need Really Big Mask to Balance Resources

- Let's make another table: **% of peak compute**
  - for **1D** tiled convolution,
  - with **TILE\_WIDTH 1024**

MASK_WIDTH	~2010 GPU with 1,000 GFLOP/s 150 GB/s	~2020 GPU with 5,000 GFLOP/s 192 GB/s
5	37%	9.6%
9	67%	17%
15	100%	28%
55	100%	100%

# Need Really Big Mask to Balance Resources

- And one more: **% of peak compute**
  - for **2D** tiled convolution,
  - with **TILE\_WIDTH 32×32**

MASK_WIDTH	~2010 GPU with 1,000 GFLOP/s 150 GB/s	~2020 GPU with 5,000 GFLOP/s 192 GB/s
3	60%	15%
5	100%	37%
7	100%	67%
9	100%	almost 100%

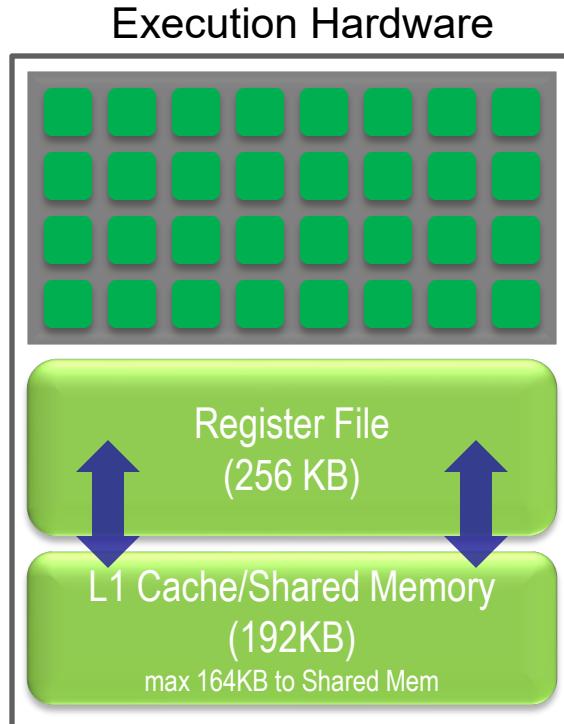
# Food for Thought

- Ratios are different for tiles on boundaries
- More importantly,
  - Each thread loads 4B to shared memory
  - 2,048 threads load only 8kB
  - Shared memory is usually 64kB or larger
  - What can one do with the rest?

**Improved approach left as homework**

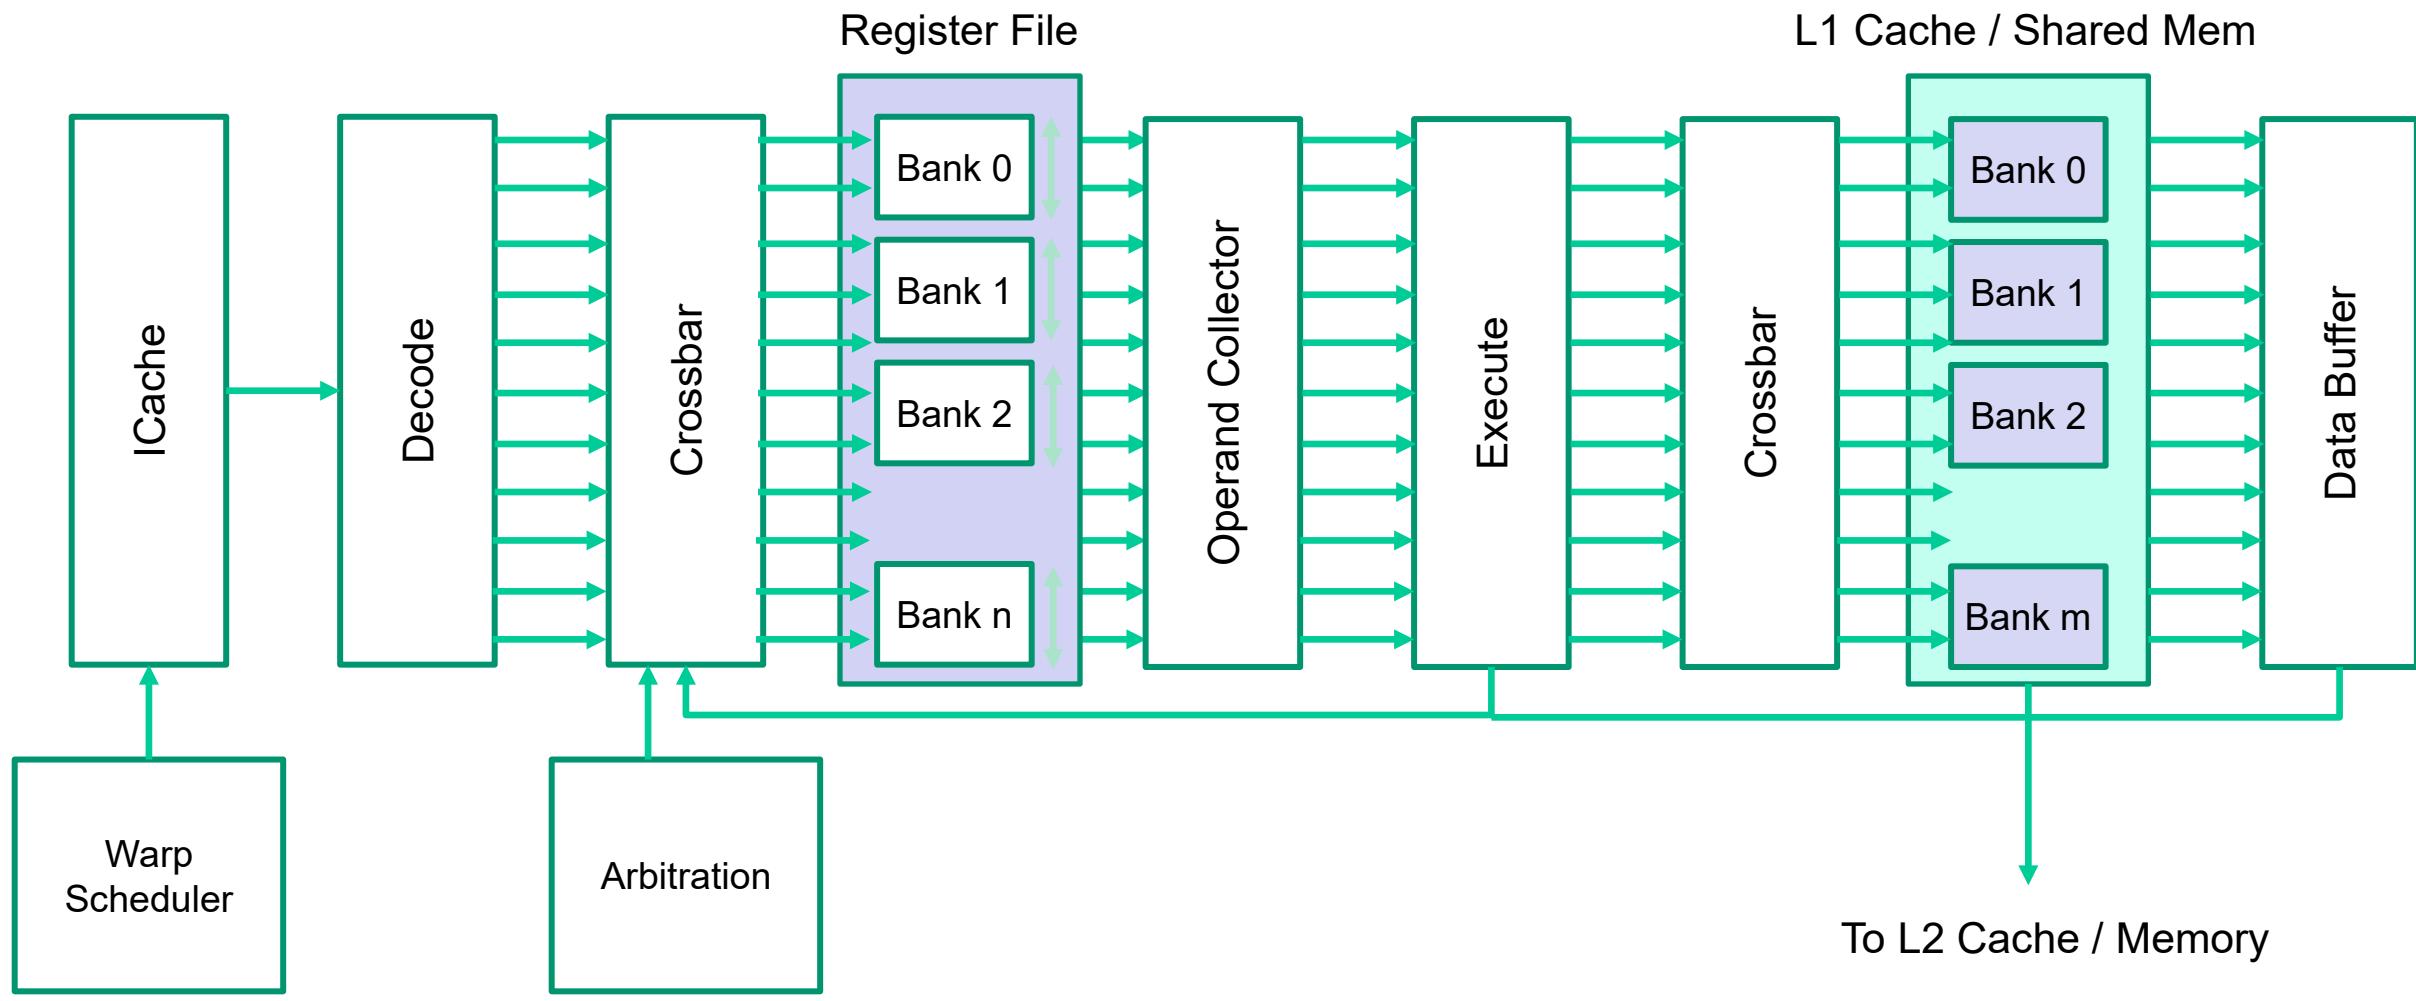
(For example, can raise MW=7 from 67% to 81%)

# Ampere SM Memory Architecture



- **Memory access time**
  - registers (~1 cycle)
  - shared memory (~5 cycles)
  - cache/constant memory (~5 cycles)
  - global memory (~500 cycles)
- **Number of thread blocks mapped to an SM is limited by**
  - **Total number of threads supported by an SM**
  - **Amount of shared memory available on SM**

# Overall Data Parallel Pipeline



# Memory Hierarchy Considerations

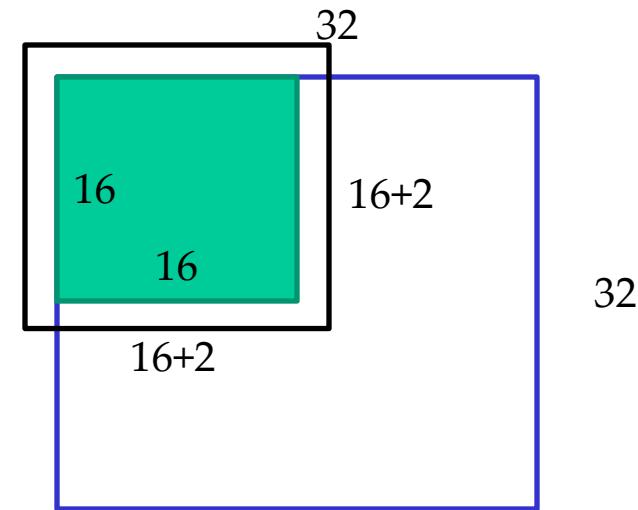
- Register file is highly banked
  - But we can have bank conflicts that cause pipeline stalls
- Shared memory is highly banked
  - But we can have bank conflicts that cause pipeline stalls
- Global memory has multiple channels, banks, pages
  - Relies on bursting
  - Coalescing is important. Need programmer involvement.
- L1 Cache is non-coherent



**ANY MORE QUESTIONS?  
READ CHAPTER 7**

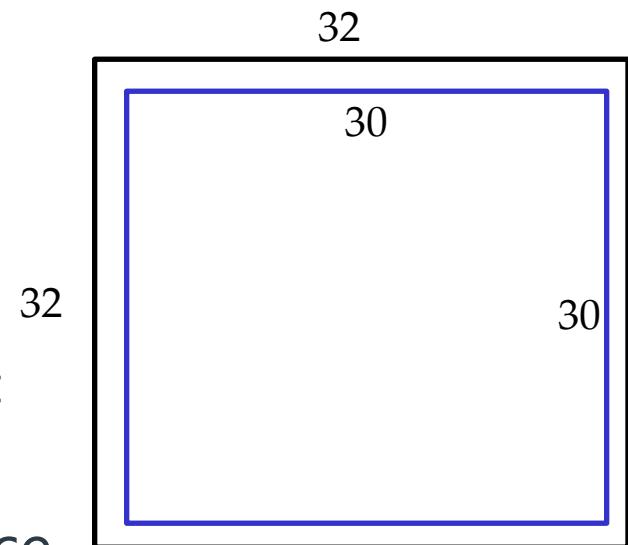
# Problem Solving

- Q: Consider 2D tiled convolution with mask width of 5x5 and output tile width of 16x16 applied to an input image of size 32x32. Assume we use Strategy 1. The use of shared memory reduces the number of global memory accesses. What is the reduction in global memory accesses for thread block (0,0)?
- A:
  - Count # of global memory accesses:  
row 0:  $9+12+15*14$   
row 1:  $12+16+20*14$   
other 14 rows:  $15+20+25*14$
  - Count # of input elements to be used in the calculation:  $18*18$



# Problem Solving

- Q: For a tiled 2D convolution kernel with 30x30 output tiles and 3x3 mask (and thus 32x32 input tile), how many warps in each thread block have control divergence? (Assume strategy 2: block size covers input tile.)
- A:
  - thread block size is 32x32
  - All 32x32 threads participate in data load
  - What does each warp do?
    - Warp 0: load 0s and do not compute
    - Warps 1-30: 30 threads compute, and 2 threads do not
    - Warp 31: load 0s and do not compute
  - Thus, only 2 warps do not have control divergence



# Problem Solving

- Q: The A40 GPU has a peak FP32 performance of 37.4 TFLOPS and 48 GB of GDDR6 memory with a memory bandwidth of 696 GB/s. These GPUs do not support FP64. What should be the byte-to-FLOP ratio for this GPU to fully utilize its compute capability? How much of data reuse is needed to make full use of the compute resources?
- A:
  - $696 \text{ GB/s} / 37.4 \text{ TFLOPS} = 0.019 \text{ B/FLOP}$ 
    - Another way to look at it: for each byte loaded from global memory, we need to use it in 53.7 floating point operations
  - This depends on the B/FLOP ratio ( $R$ ) of the GPU kernel
    - $100 / (696 \text{ GB/s} / R \text{ B/FLOP})$