



ECE408/CS483/CSE408 Fall 2024

# Applied Parallel Programming

## Lecture 12: Computation in Deep Neural Networks

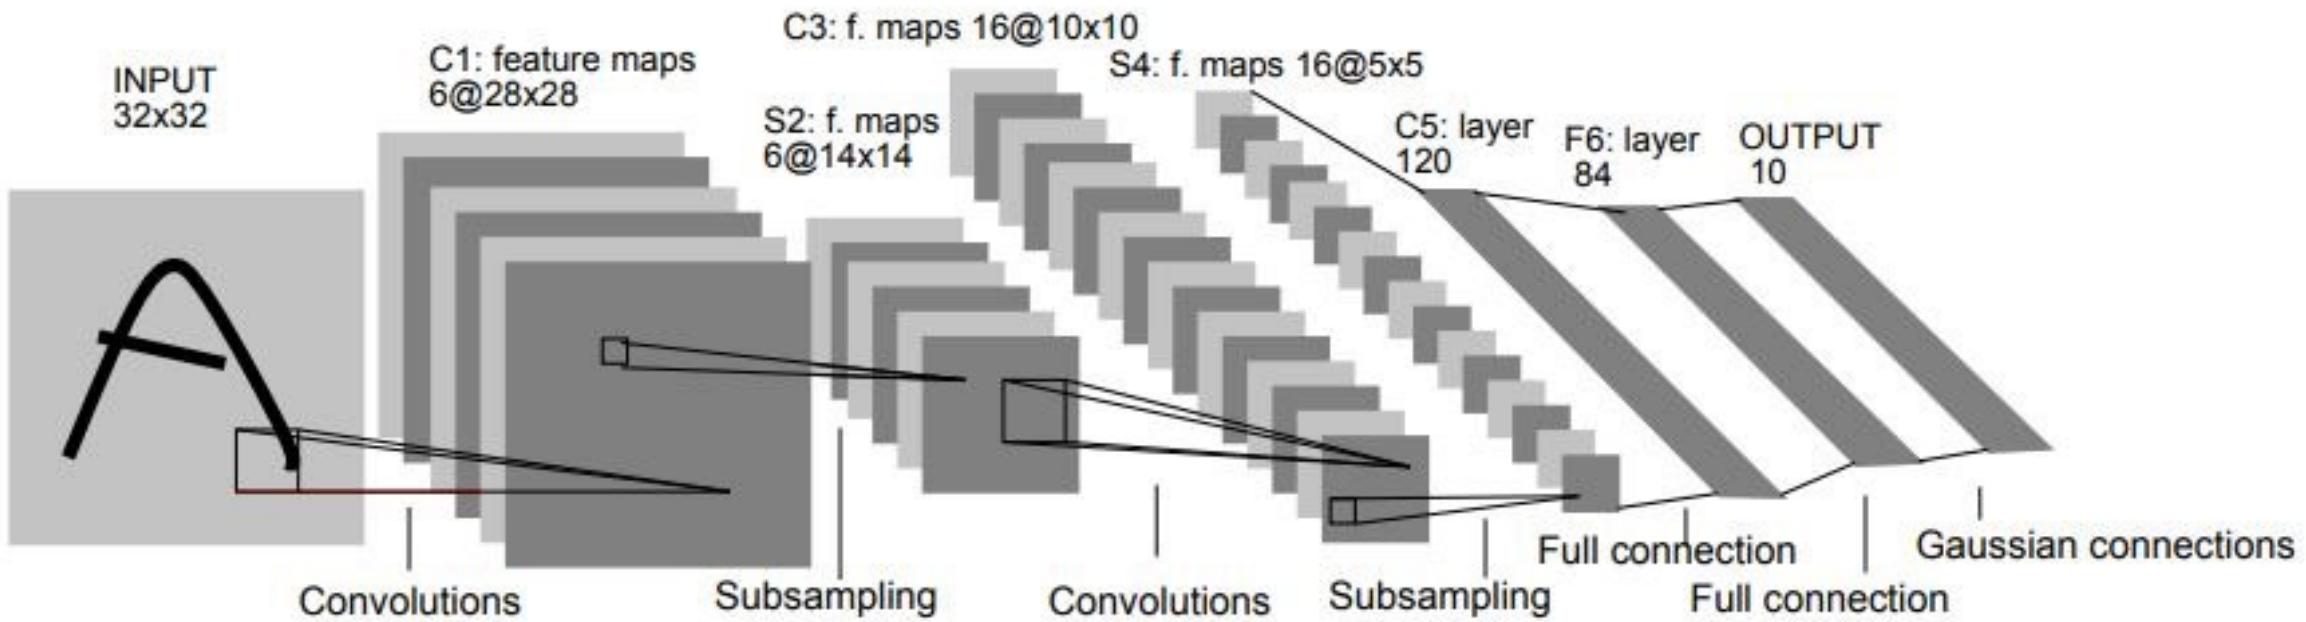
# Course Reminders

- Midterm 1 is on Tuesday, October 15<sup>th</sup>
  - See Canvas for details, including exact time and topics
  - Please email your lecture instructor if you have a conflict by this Monday, October. 7<sup>th</sup>
- Lab 4 due this Friday
- Project Milestone 1
  - Due on October 11<sup>th</sup>
- Next lecture is pre-recorded, please watch it on your own time.
  - The material covered in this lecture will be assessed in the upcoming midterm exam.

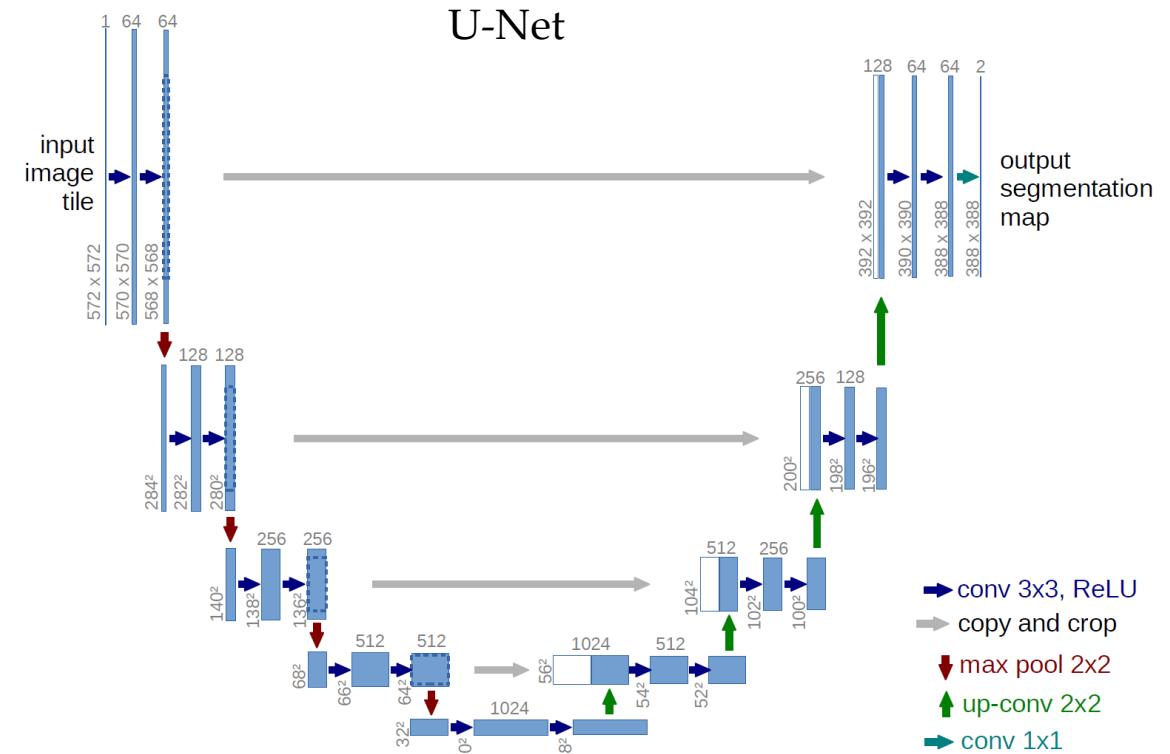
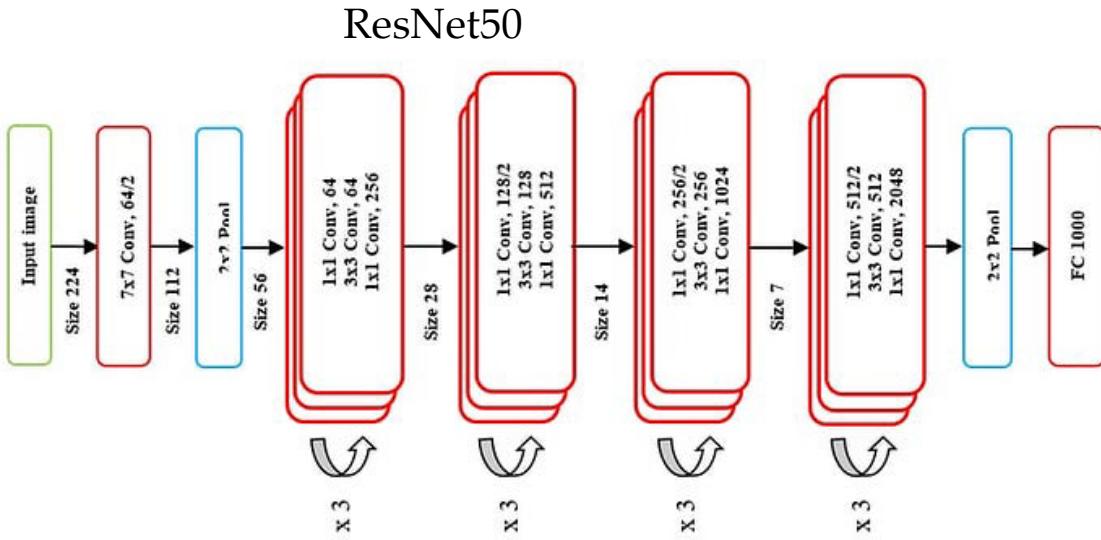
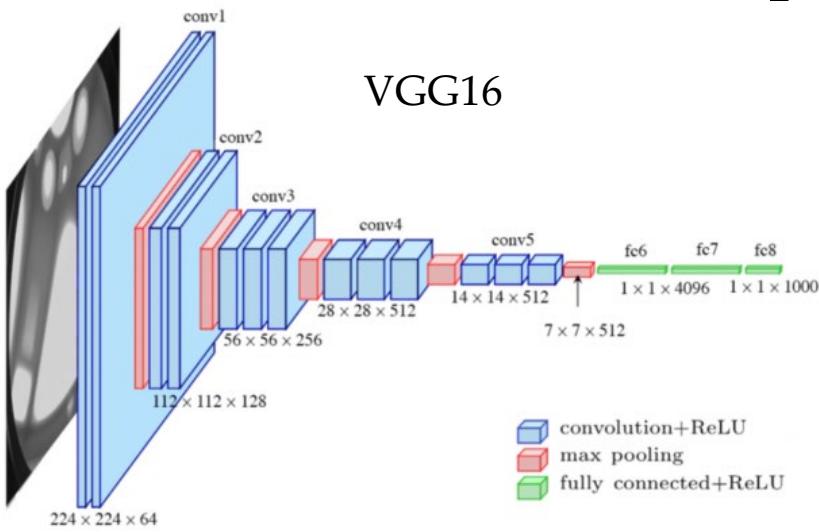
# Objective

- To learn to implement the different types of layers in a Convolutional Neural Network (CNN)
- Brief intro to Project Milestone 1

# LeNet-5:CNN for hand-written digit recognition



# Many Types of CNNs



# Anatomy of a Convolution Layer

## Input features

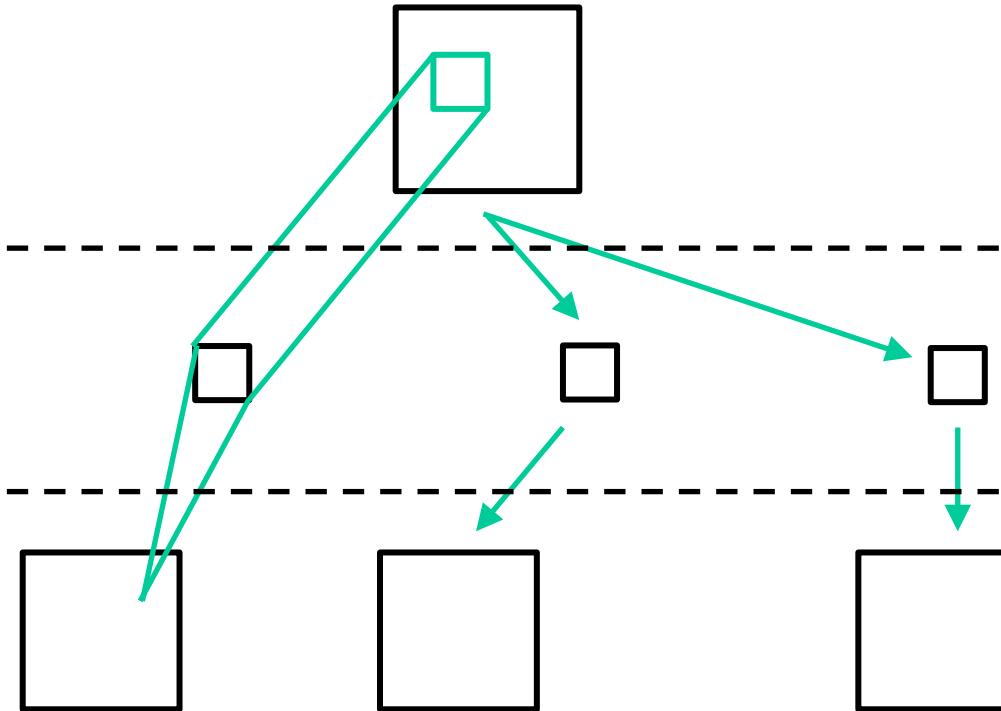
- A inputs each  $N_1 \times N_2$

## Convolution Layer

- B convolution kernels each  $K_1 \times K_2$

## Output Features (total of B)

- $A \times B$  outputs each  
 $(N_1 - K_1 + 1) \times (N_2 - K_2 + 1)$



# Notion of a Channel in Input Layer

Some Set of Input Features are Related

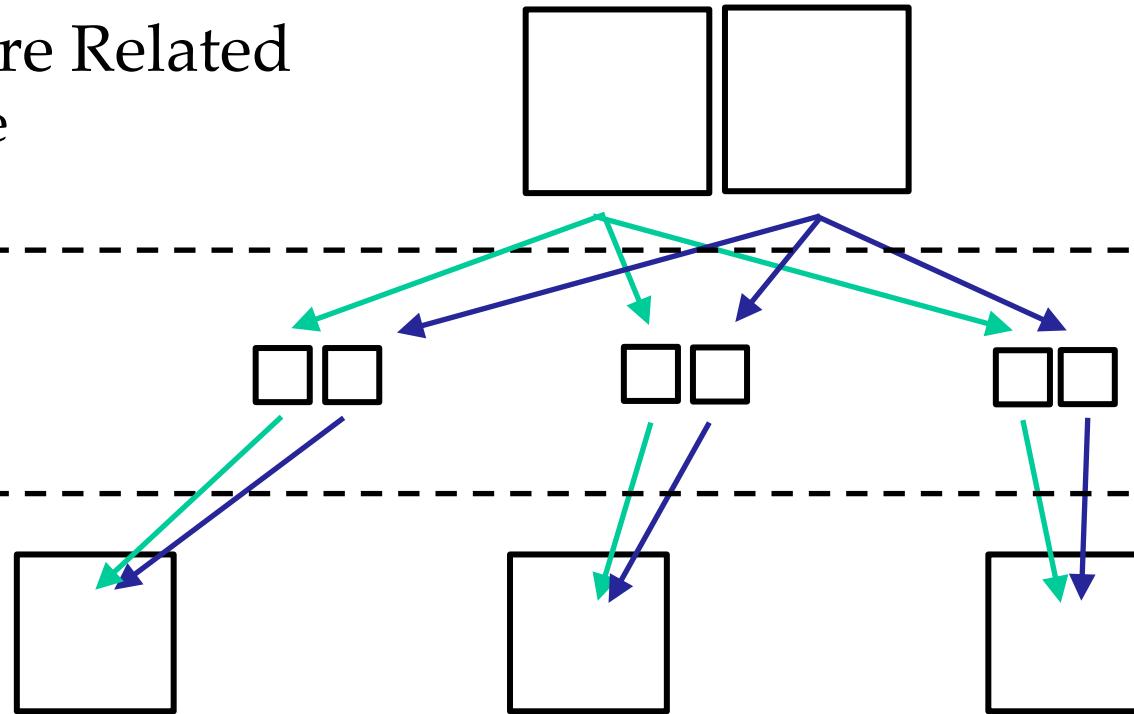
- For example: Red, Green, Blue

Convolution Layer

- Different kernels per channel

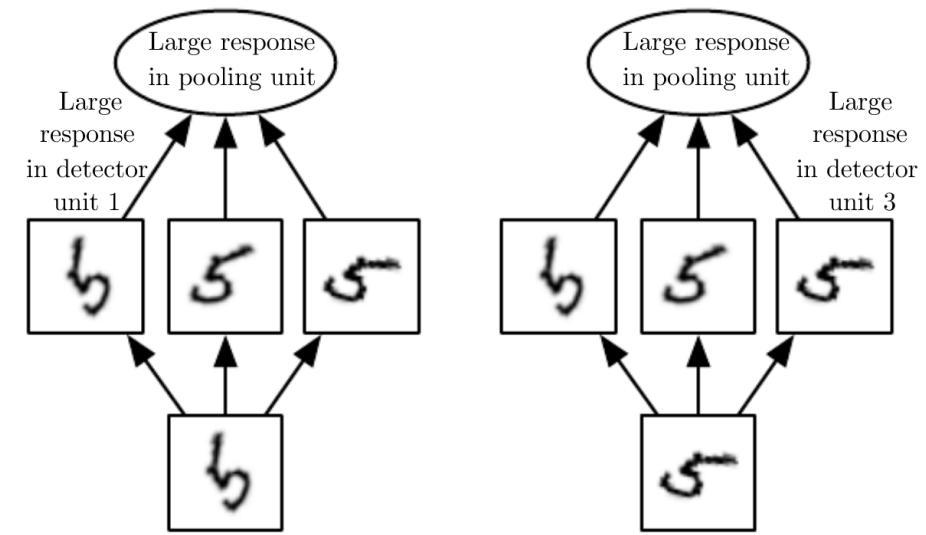
Output Features

- Channels combine per output feature

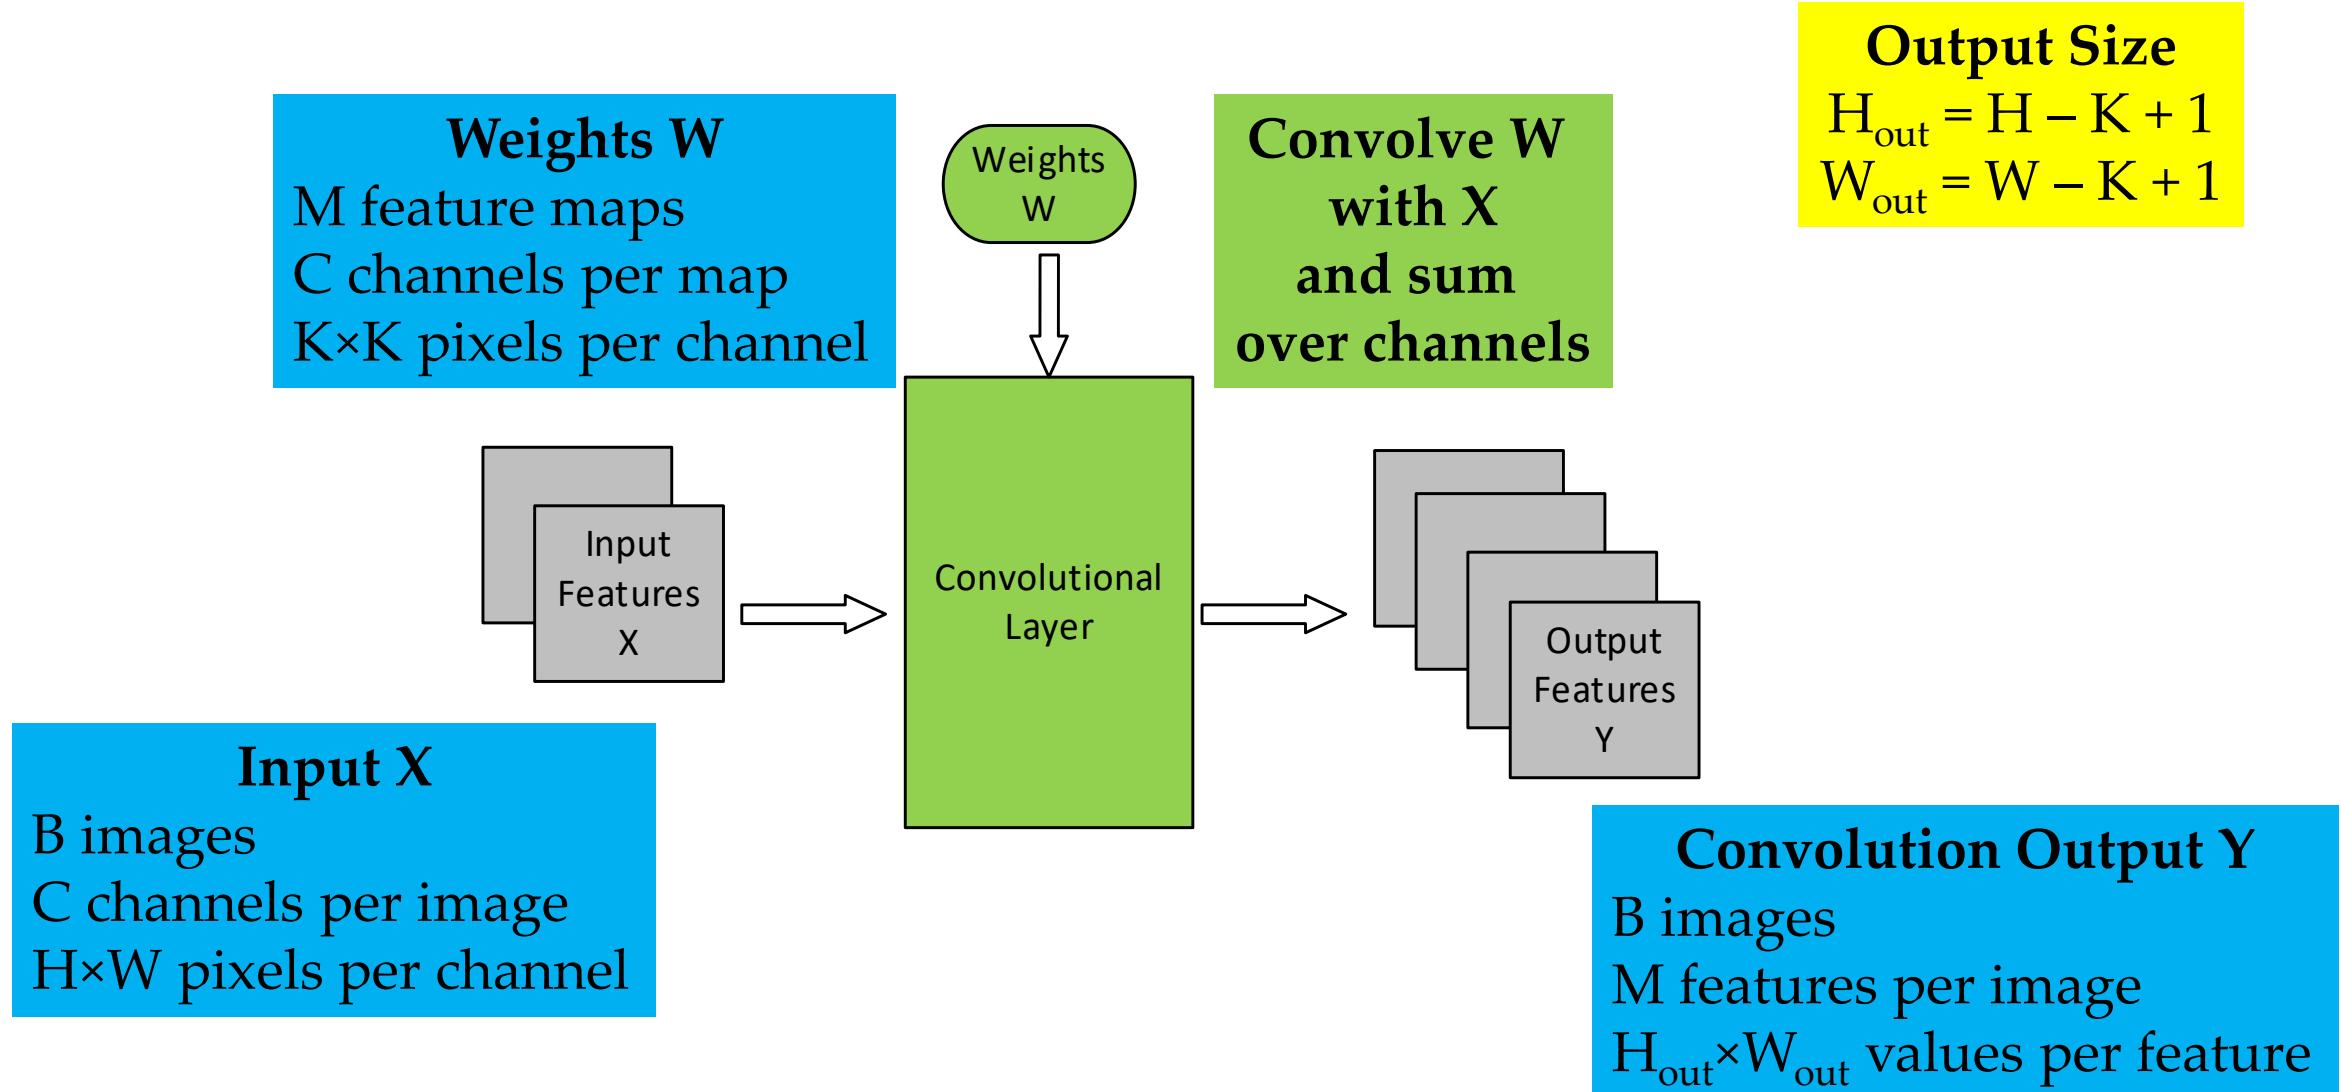


# 2-D Pooling (Subsampling)

- A subsampling layer
  - Sometimes with bias and non-linearity built in
- Common types
  - max, average,  $L^2$  norm, weighted average
- Helps make representation invariant to size scaling and small translations in the input



# Forward Propagation

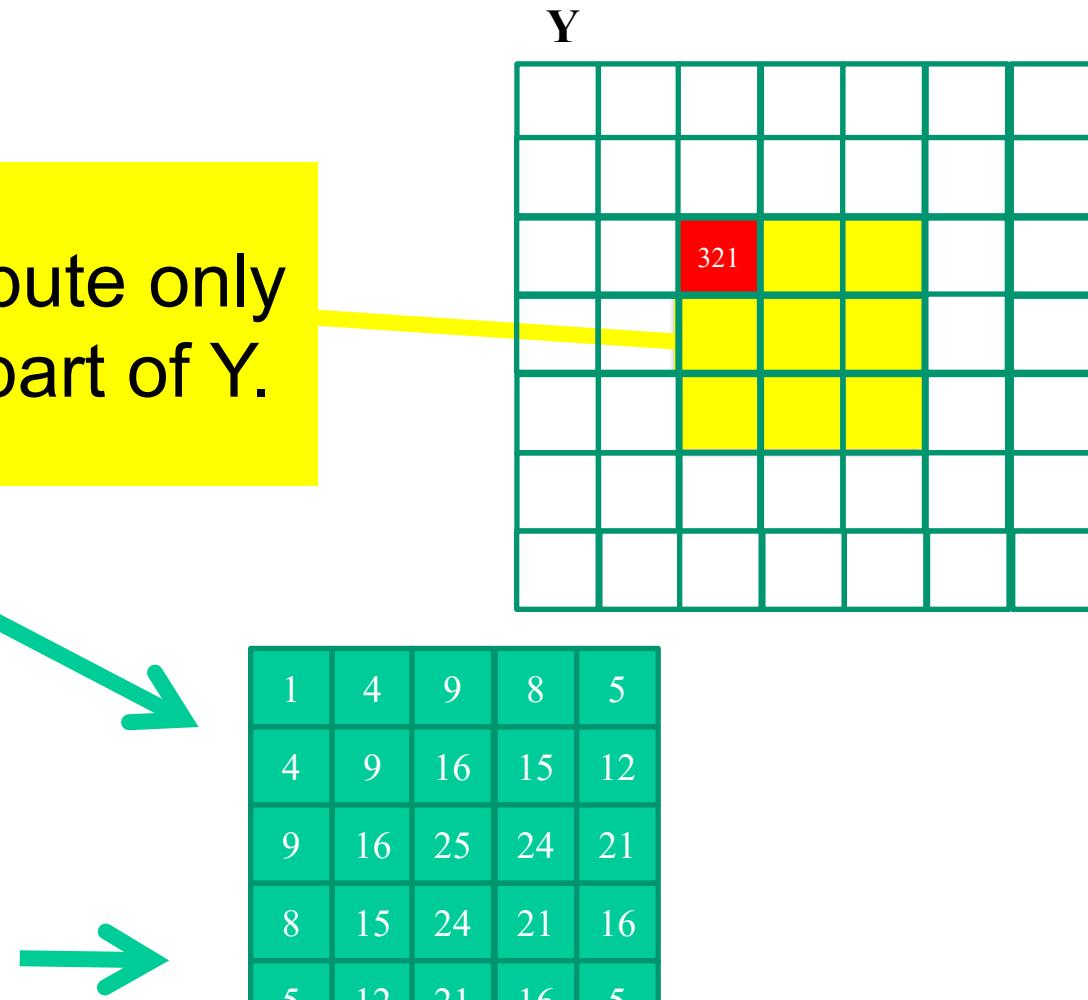


# Outputs Must Use Full Mask/Kernel

X	1	2	3	4	5	6	7
2	1	3	4	5	6	7	8
3	2	4	5	6	7	8	9
4	3	5	6	7	8	5	6
5	4	6	7	8	5	6	7
6	5	7	8	9	0	1	2
7	6	8	9	0	1	2	3

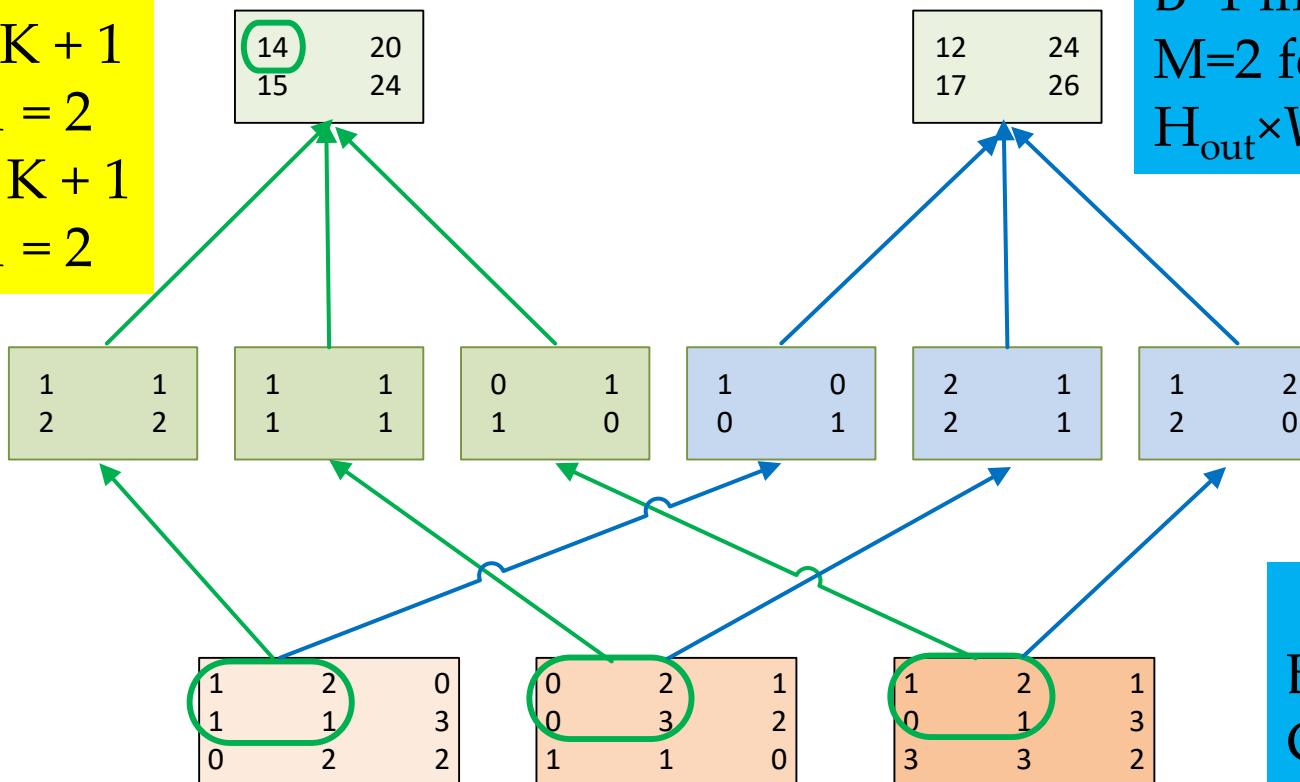
W	1	2	3	2	1
2	1	2	3	4	3
3	2	3	4	5	4
4	3	4	5	4	3
5	2	3	4	3	2
6	1	2	3	2	1

Compute only  
this part of Y.



# Example of the Forward Path of a Convolution Layer

**Output Size**  
 $H_{out} = H - K + 1$   
 $= 3 - 2 + 1 = 2$   
 $W_{out} = W - K + 1$   
 $= 3 - 2 + 1 = 2$



**Convolution Output Y**  
B=1 image  
M=2 features per image  
 $H_{out} \times W_{out} = 2 \times 2$  values per feature

**Weights W**  
M=2 feature maps  
C=3 channels per map  
 $K \times K = 2 \times 2$  pixels per channel

**Input X**  
B=1 image  
C=3 channels  
 $H \times W = 3 \times 3$  pixels per channel

# Sequential Code: Forward Convolutional Layer

```
void convLayer_forward(int B, int M, int C, int H, int W, int K, float* X, float* W, float* Y)
{
    int H_out = H - K + 1;                                // calculate H_out, W_out
    int W_out = W - K + 1;

    for (int b = 0; b < B; ++b)                          // for each image
        for(int m = 0; m < M; m++)                      // for each output feature map
            for(int h = 0; h < H_out; h++)                // for each output value (two loops)
                for(int w = 0; w < W_out; w++) {
                    Y[b, m, h, w] = 0.0f;                  // initialize sum to 0
                    for(int c = 0; c < C; c++)           // sum over all input channels
                        for(int p = 0; p < K; p++)       // KxK filter
                            for(int q = 0; q < K; q++)
                                Y[b, m, h, w] += X[b, c, h + p, w + q] * W[m, c, p, q];
                }
            }
}
```

# A Small Convolution Layer Example

$X[b, 1, \dots]$  input channel  
 $W[0,1,\dots]$   
 $Y[b, 0, \dots]$  output map

Image  $b$  in mini batch  
 $X[b,0,\dots]$

1	2	0	1
1	1	3	2
0	2	2	0
2	1	0	3

1	1	1
2	2	3
2	1	0

$w[0,0,\dots]$

$X[b,1,\dots]$

0	2	1	0
0	3	2	1
1	1	0	2
2	1	0	3

1	2	3
1	1	0
3	0	1

$w[0,1,\dots]$

0	?
?	?
?	?

$y[b,0,\dots]$

$X[b,2,\dots]$

1	2	1	0
0	1	3	2
3	3	2	0
1	3	2	0

0	1	1
1	0	2
1	2	1

$w[0,2,\dots]$

# A Small Convolution Layer Example

$c = 0$

$x[b,0,\_,\_]$

1	2	0	1
1	1	3	2
0	2	2	0
2	1	0	3

$x[b,1,\_,\_]$

0	2	1	0
0	3	2	1
1	1	0	2
2	1	0	3

$x[b,2,\_,\_]$

1	2	1	0
0	1	3	2
3	3	2	0
1	3	2	0

$w[0,0,\_,\_]$

1	1	1
2	2	3
2	1	0

$$3+13+2$$

$y[b,0,\_,\_]$

18	?
?	?

$w[0,1,\_,\_]$

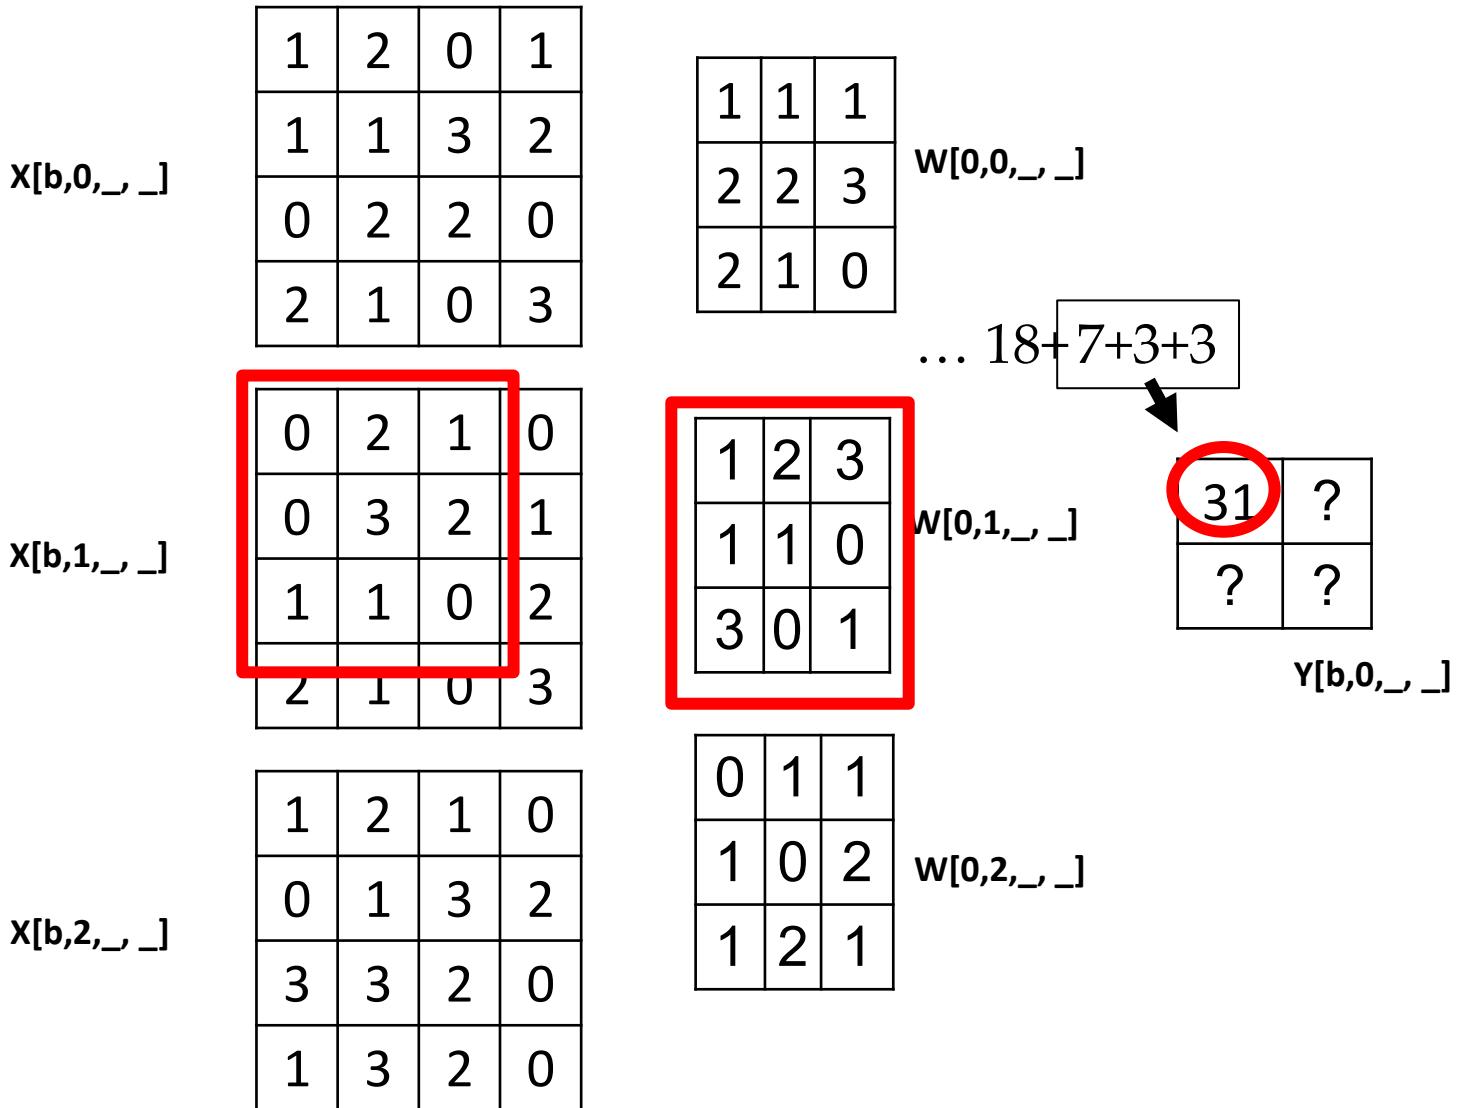
1	2	3
1	1	0
3	0	1

$w[0,2,\_,\_]$

0	1	1
1	0	2
1	2	1

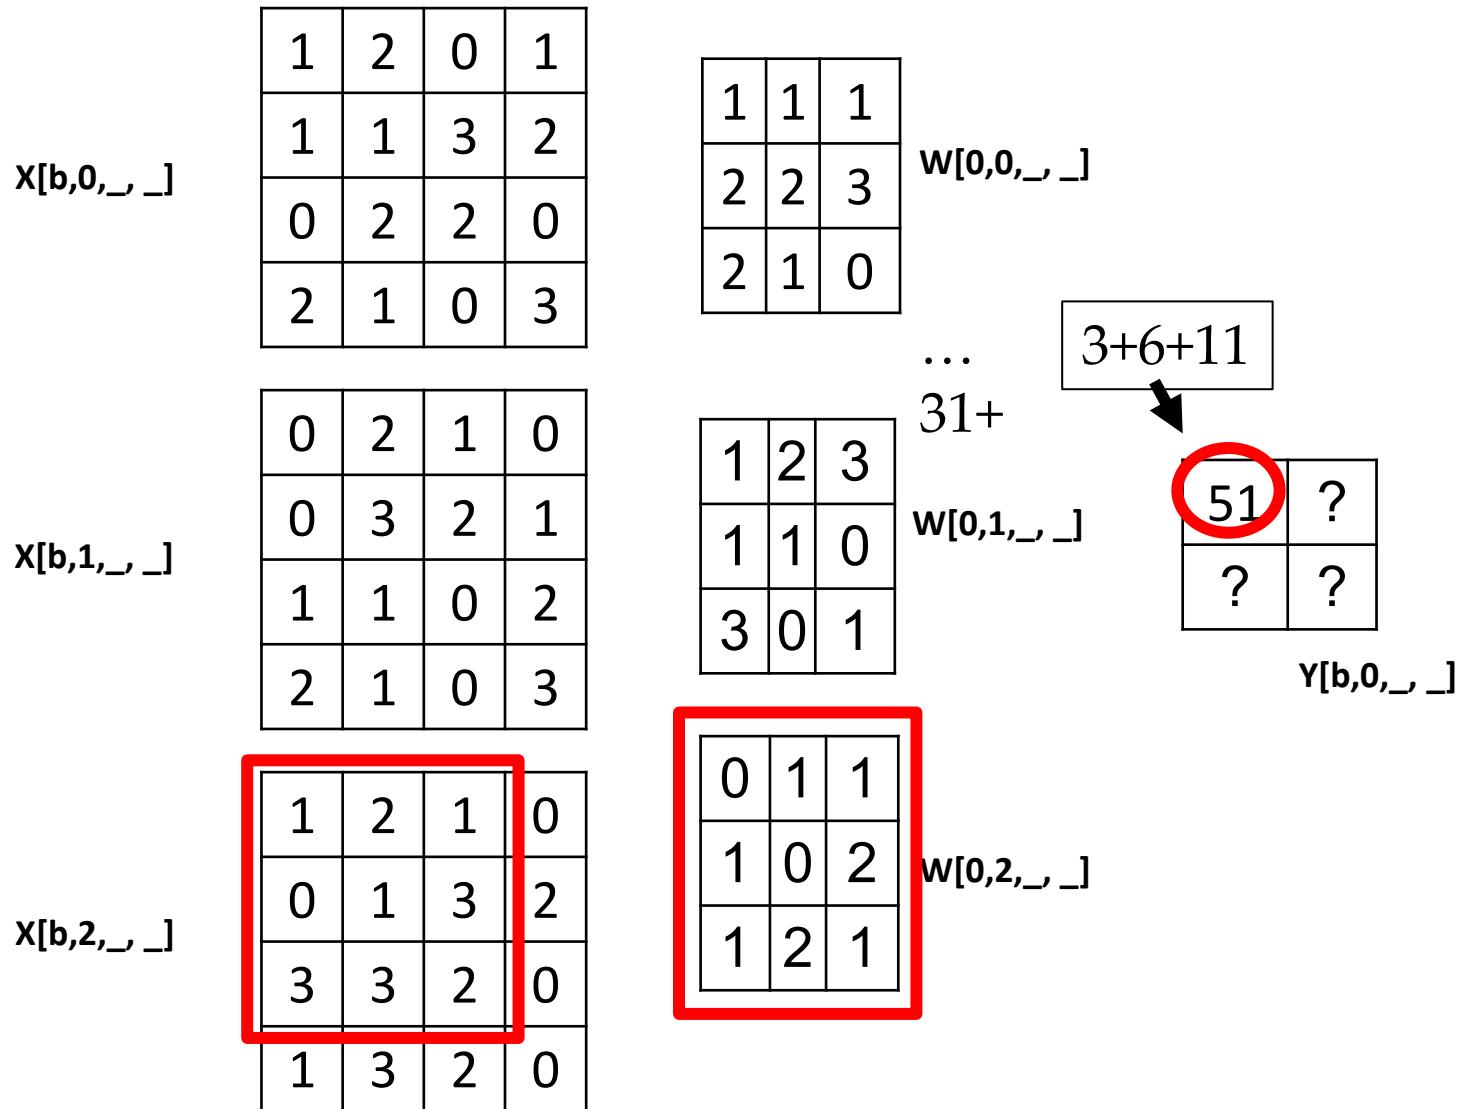
# A Small Convolution Layer Example

$c = 1$



# A Small Convolution Layer Example

$c = 2$



# Parallelism in a Convolution Layer

**Output feature maps** can be calculated in parallel

- Usually a small number, not sufficient to fully utilize a GPU

All **output** feature map **pixels** can be calculated in parallel

- All rows can be done in parallel
- All pixels in each row can be done in parallel
- Large number but diminishes as we go into deeper layers

All **input feature maps** can be processed in parallel,  
but need atomic operation or tree reduction (we'll learn later)

**Different layers may demand different strategies.**

# Subsampling (Pooling) by Scale N

**Convolution Output Y**  
B images  
M features per image  
 $H_{out} \times W_{out}$  values per feature

Average over  $N \times N$  blocks,  
then calculate sigmoid

**Output Size**  
 $H_{S(N)} = \text{floor}(H_{out} / N)$   
 $W_{S(N)} = \text{floor}(W_{out} / N)$

**Subsampling/Pooling Output S**  
B images  
M features per image  
 $H_{S(N)} \times W_{S(N)}$  values per feature

# Sequential Code: Forward Pooling Layer

```
void poolingLayer_forward(int B, int M, int H_out, int W_out, int N, float* Y, float* S)
{
    for (int b = 0; b < B; ++b)                      // for each image
        for (int m = 0; m < M; ++m)                  // for each output feature map
            for (int x = 0; x < H_out/N; ++x)          // for each output value (two loops)
                for (int y = 0; y < W_out/N; ++y) {
                    float acc = 0.0f                      // initialize sum to 0
                    for (int p = 0; p < N; ++p)           // loop over NxN block of Y (two loops)
                        for (int q = 0; q < N; ++q)
                            acc += Y[b, m, N*x + p, N*y + q];
                    acc /= N * N;                      // calculate average over block
                    S[b, m, x, y] = sigmoid(acc + bias[m]) // bias, non-linearity
                }
}
```

# Kernel Implementation of Subsampling Layer

- Straightforward mapping from grid to subsampled output feature map pixels
- in GPU kernel,
  - need to manipulate index mapping
  - for accessing the output feature map pixels
  - of the previous convolution layer.
- Often merged into the previous convolution layer to save memory bandwidth

# Design of a Basic Kernel

- Each block computes
  - a tile of output pixels for one feature
  - `TILE_WIDTH` pixels in each dimension
- Grid's X dimension maps to M output feature maps
- Grid's Y dimension maps to the tiles in the output feature maps (linearized order).
- (Grid's Z dimension is used for images in batch, which we omit from slides.)

tiles covering an output feature map, marked with linearized indices

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19

# A Small Example

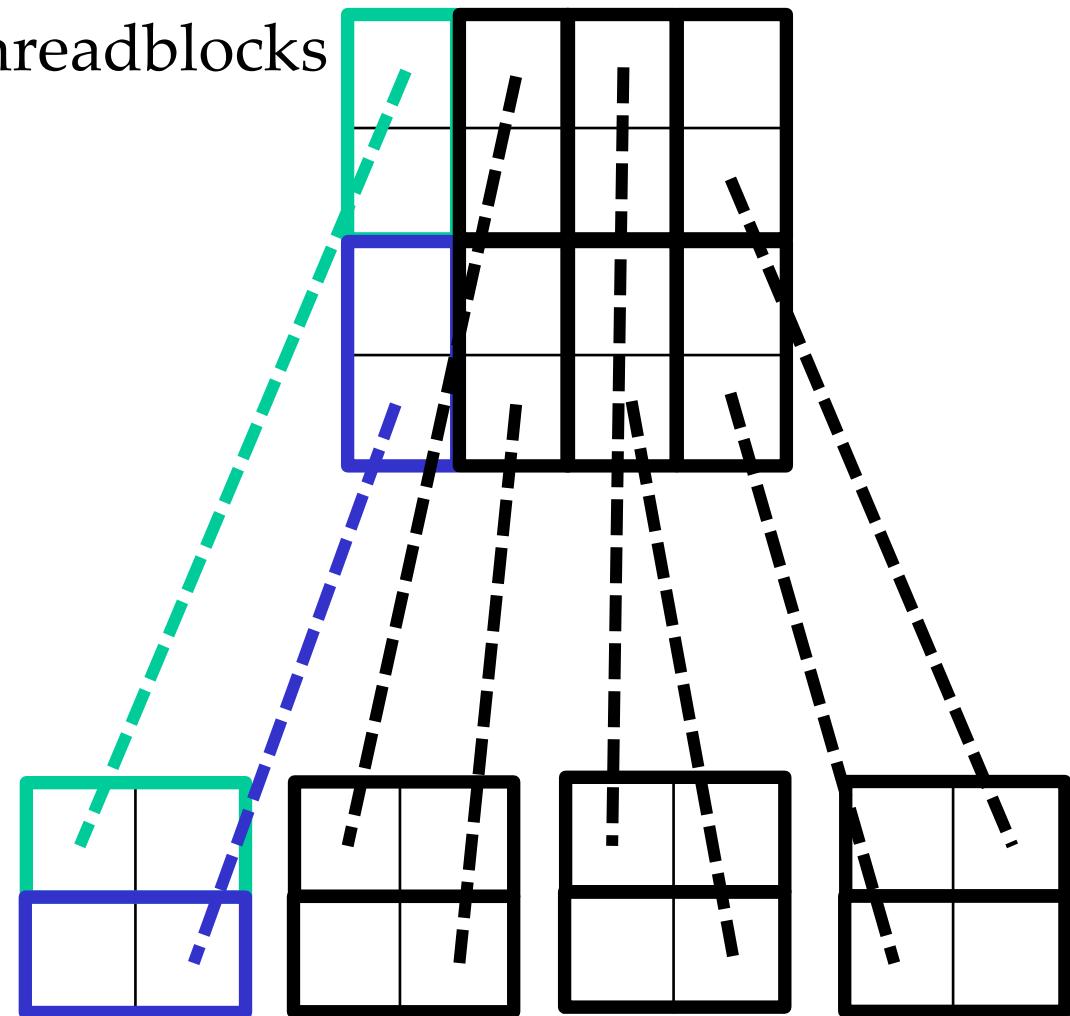
Assume

- **M = 4** (4 output feature maps),
- thus 4 blocks in the X dimension, and
- **W\_out = H\_out = 8** (8x8 output features).

If **TILE\_WIDTH = 4**,  
we also need 4 blocks in the Y  
dimension:

- for each output feature,
- top two blocks in each column  
calculates the top row of tiles, and
- bottom two calculate the bottom row.

CUDA Grid and  
Threadblocks



Output Feature Maps and Tiles <sup>22</sup>

# Host Code for a Basic Kernel: CUDA Grid

Consider an output feature map:

- width is **W\_out**, and
- height is **H\_out**.
- Assume these are multiples of **TILE\_WIDTH**.

0	1	2	3	4
5	6	7	8	9
10	11	12	13	14
15	16	17	18	19

Let **X\_grid** be the number of blocks needed in X dim (5 above).

Let **Y\_grid** be the number of blocks needed in Y dim (4 above).

# Host Code for a Basic Kernel: CUDA Grid

(Assuming W\_out and H\_out are multiples of TILE\_WIDTH.)

```
#define TILE_WIDTH 16          // We will use 4 for small examples.  
W_grid = W_out/TILE_WIDTH;   // number of horizontal tiles per output map  
H_grid = H_out/TILE_WIDTH;   // number of vertical tiles per output map  
Y = H_grid * W_grid;  
  
dim3 blockDim(TILE_WIDTH, TILE_WIDTH, 1); // output tile for untiled code  
dim3 gridDim(M, Y, 1);  
  
ConvLayerForward_Kernel<<< gridDim, blockDim >>>(...);
```

# Sequential Code: Forward Convolutional Layer

```
void convLayer_forward(int B, int M, int C, int H, int W, int K, float* X, float* W, float* Y)
{
    int H_out = H - K + 1;                                // calculate H_out, W_out
    int W_out = W - K + 1;

    for (int b = 0; b < B; ++b)                          // for each image
        for (int m = 0; m < M; m++)                      // for each output feature map
            for (int h = 0; h < H_out; h++)                // for each output value (two loops)
                for (int w = 0; w < W_out; w++) {
                    Y[b, m, h, w] = 0.0f;                  // initialize sum to 0
                    for (int c = 0; c < C; c++)           // sum over all input channels
                        for (int p = 0; p < K; p++)       // KxK filter
                            for (int q = 0; q < K; q++)
                                Y[b, m, h, w] += X[b, c, h + p, w + q] * W[m, c, p, q];
                }
}
```

Computed  
by the grid

Computed by a thread

# Partial Kernel Code for a Convolution Layer

```
__global__ void ConvLayerForward_Basic_Kernel
    (int C, int W_grid, int K, float* X, float* W, float* Y)
{
    int m = blockIdx.x;
    int h = (blockIdx.y / W_grid) * TILE_WIDTH + threadIdx.y;
    int w = (blockIdx.y % W_grid) * TILE_WIDTH + threadIdx.x;
    float acc = 0.0f;
    for (int c = 0; c < C; c++) {                      // sum over all input channels
        for (int p = 0; p < K; p++) {                  // loop over KxK filter
            for (int q = 0; q < K; q++)
                acc += X[c, h + p, w + q] * W[m, c, p, q];
        }
        Y[m, h, w] = acc;
    }
}
```

# Some Observations

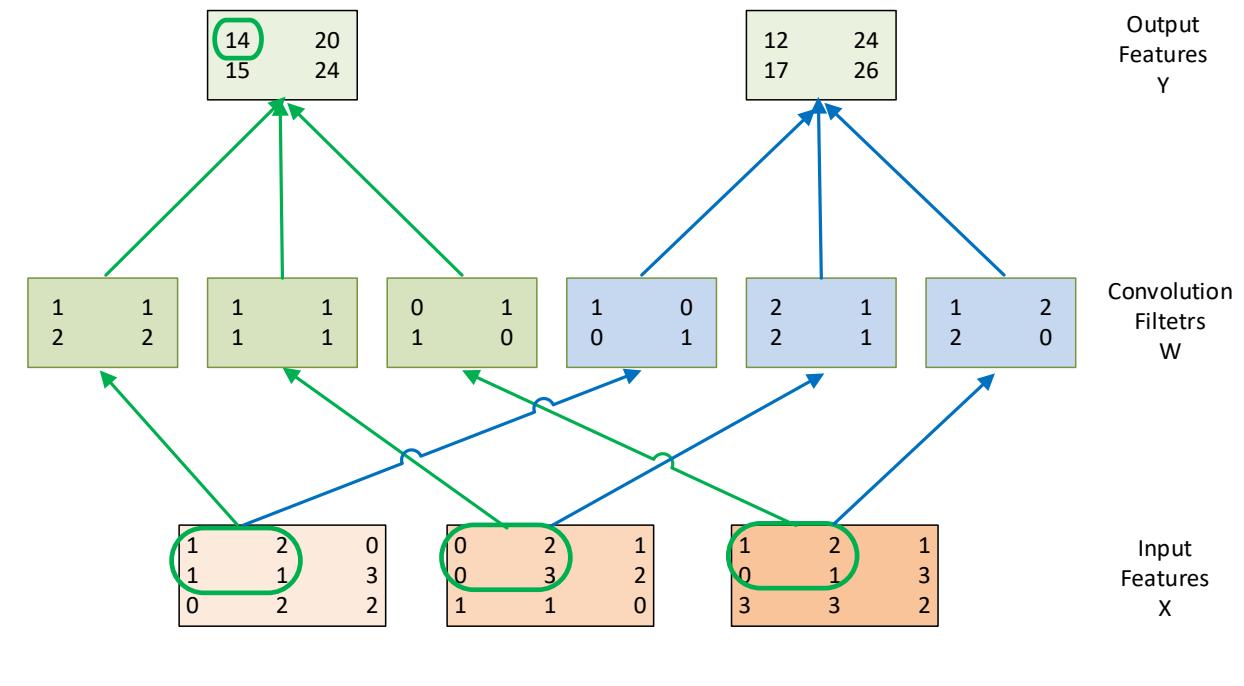
## Enough parallelism

- if the total number of pixels
- across all output feature maps is large
- (often the case for CNN layers)

## Each input tile

- loaded M times (number of output features), so
- **not efficient in global memory bandwidth,**
- but block scheduling in X dimension should give cache benefits.

# Implementing a Convolution Layer with Matrix Multiplication



$$\begin{array}{c} \begin{array}{|c|c|c|} \hline 1 & 1 & 2 & 2 \\ \hline 1 & 0 & 0 & 1 \\ \hline \end{array} & \begin{array}{|c|c|c|} \hline 1 & 1 & 1 & 1 \\ \hline 2 & 1 & 2 & 1 \\ \hline \end{array} & \begin{array}{|c|c|c|} \hline 0 & 1 & 1 & 0 \\ \hline 1 & 2 & 2 & 0 \\ \hline \end{array} \\ \hline \end{array} * \begin{array}{|c|c|c|} \hline 1 & 2 & 1 & 1 \\ \hline 1 & 1 & 0 & 2 \\ \hline 1 & 3 & 2 & 2 \\ \hline 0 & 2 & 0 & 3 \\ \hline 2 & 1 & 3 & 2 \\ \hline 0 & 3 & 1 & 1 \\ \hline 3 & 2 & 1 & 0 \\ \hline 1 & 2 & 1 & 1 \\ \hline 2 & 1 & 0 & 3 \\ \hline 0 & 1 & 3 & 3 \\ \hline 1 & 3 & 3 & 2 \\ \hline \end{array} = \begin{array}{|c|c|} \hline 14 & 20 \\ \hline 12 & 24 \\ \hline 15 & 24 \\ \hline 17 & 26 \\ \hline \end{array}$$

Convolution Filters  $W'$

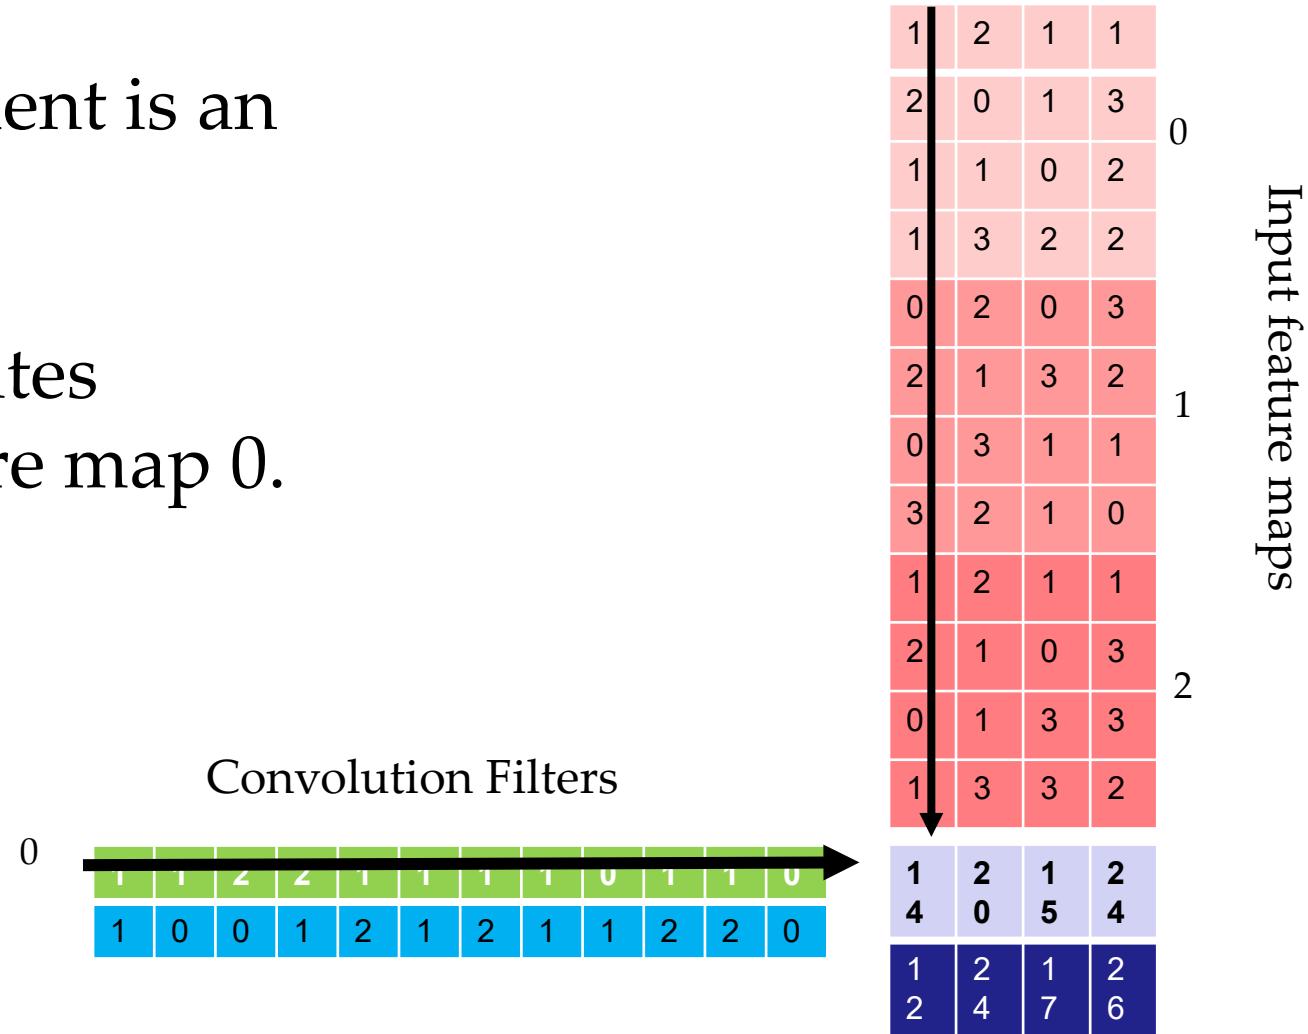
Output Features  $Y$

Input Features  $X_{\text{unrolled}}$

# Simple Matrix Multiplication

Each product matrix element is an output feature map pixel.

This inner product generates element 0 of output feature map 0.

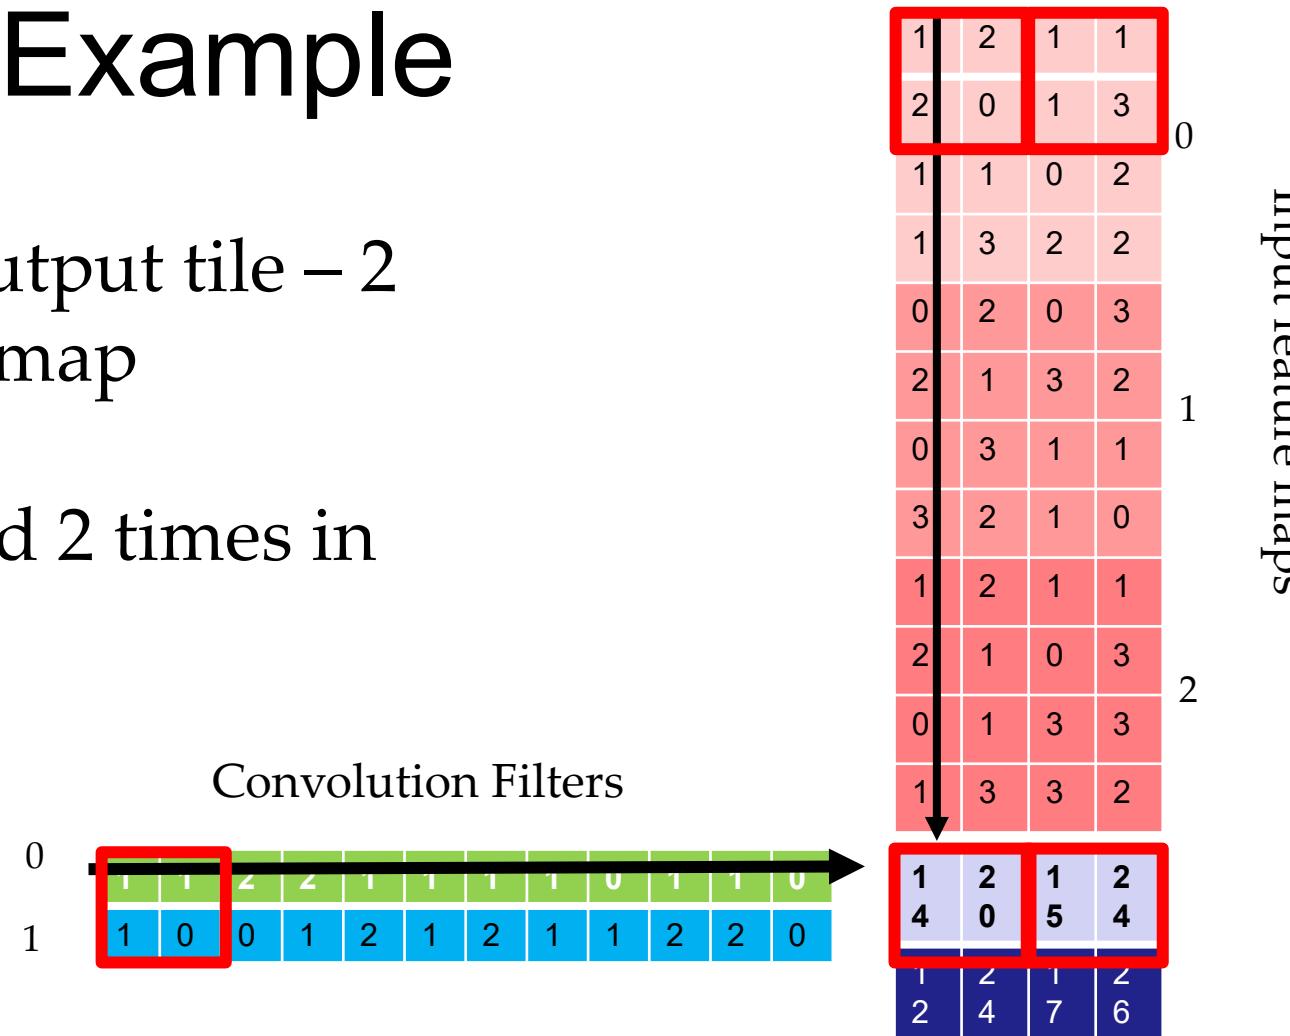


# Tiled Matrix Multiplication

## 2x2 Example

Each block calculates one output tile – 2 elements from each output map

Each input element is reused 2 times in the shared memory

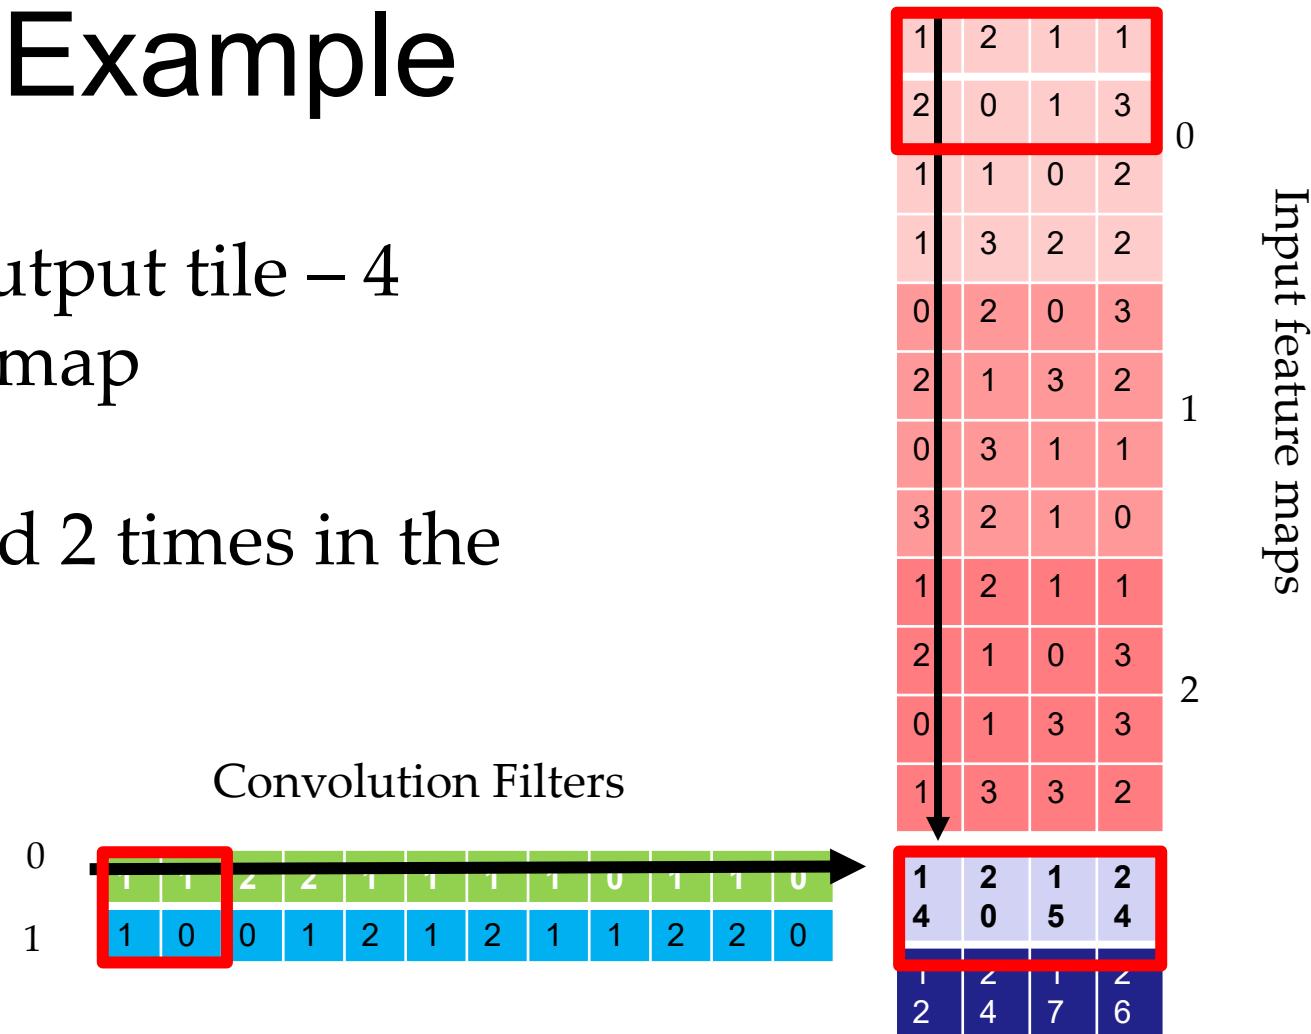


# Tiled Matrix Multiplication

## 2x4 Example

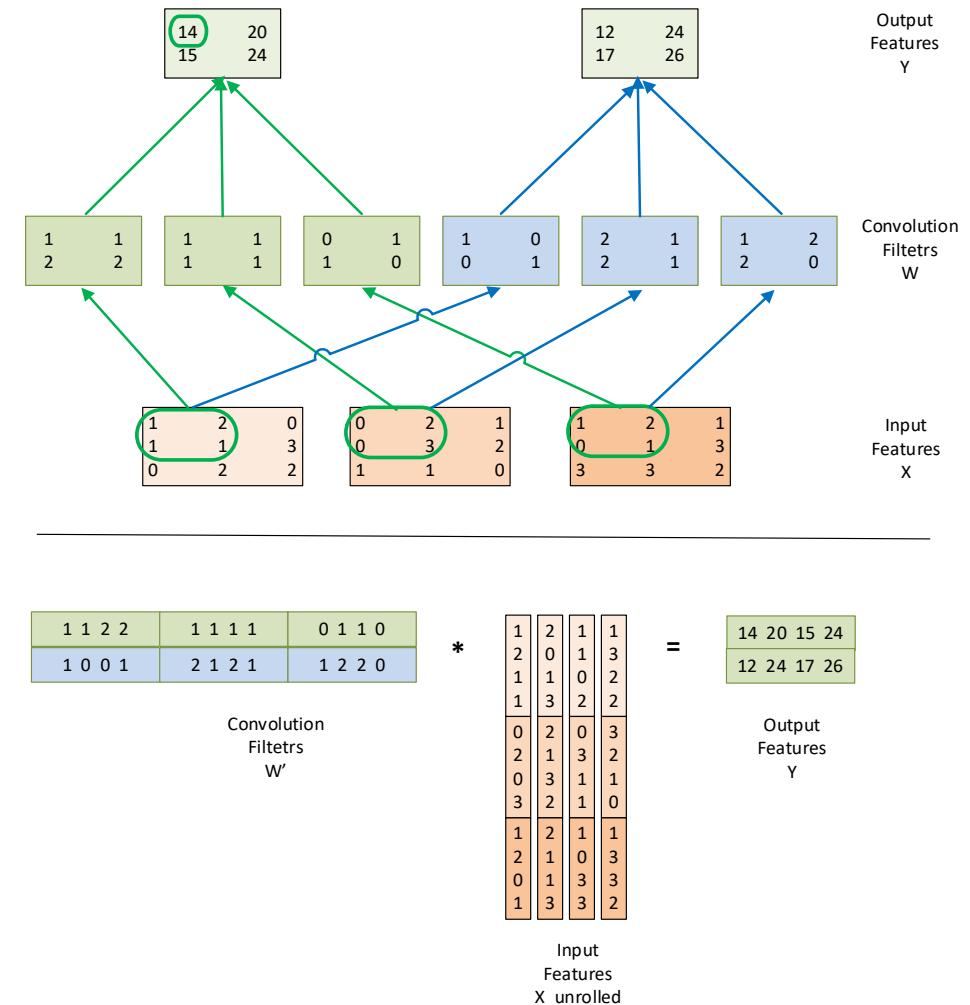
Each block calculates one output tile – 4 elements from each output map

Each input element is reused 2 times in the shared memory



# Efficiency Analysis: Total Input Replication

- Replicated input features are shared among output maps
  - There are  $H_{out} * W_{out}$  output feature map elements
  - Each requires  $K*K$  elements from the input feature maps
  - So, the total number of input element after replication is  $H_{out}*W_{out}*K*K$  times for each input feature map
  - The total number of elements in each original input feature map is  $(H_{out}+K-1) * (W_{out}+K-1)$



# Analysis of a Small Example

$$H_{out} = 2$$

$$W_{out} = 2$$

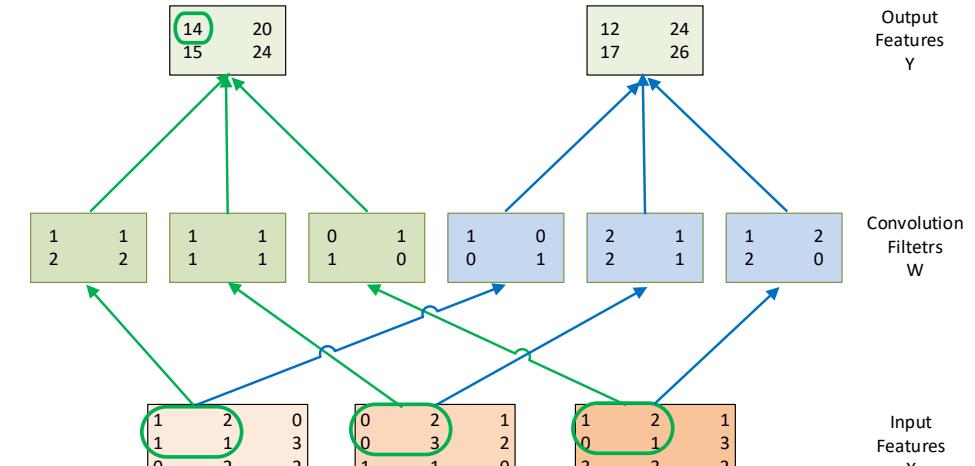
$$K = 2$$

There are 3 input maps (channels)

The total number of input elements in the replicated (“unrolled”) input matrix is  $3*2*2*2*2$

The replicating factor is

$$(3*2*2*2*2)/(3*3*3) = 1.78$$



$$\begin{array}{c|c|c} 1 & 1 & 2 \\ \hline 1 & 0 & 0 \\ \hline 2 & 1 & 1 \\ \hline 0 & 2 & 2 \end{array} \quad \begin{array}{c|c|c} 1 & 1 & 1 \\ \hline 2 & 1 & 2 \\ \hline 1 & 1 & 1 \\ \hline 0 & 1 & 0 \end{array} \quad \begin{array}{c|c|c} 0 & 1 & 1 \\ \hline 1 & 2 & 2 \\ \hline 0 & 2 & 2 \\ \hline 1 & 1 & 0 \end{array} \quad * \quad \begin{array}{c|c|c} 1 & 2 & 1 \\ \hline 2 & 0 & 1 \\ \hline 1 & 1 & 3 \\ \hline 1 & 3 & 2 \\ \hline 0 & 2 & 0 \\ \hline 2 & 1 & 3 \\ \hline 0 & 3 & 1 \\ \hline 3 & 2 & 1 \\ \hline 1 & 2 & 1 \\ \hline 2 & 1 & 0 \\ \hline 0 & 1 & 3 \\ \hline 1 & 3 & 2 \end{array} \quad = \quad \begin{array}{c|c|c} 14 & 20 & 15 \\ \hline 12 & 24 & 17 \\ \hline 26 & 24 & 26 \end{array}$$

Convolution  
Filters  
 $W'$

Input  
Features  
 $X_{\text{unrolled}}$

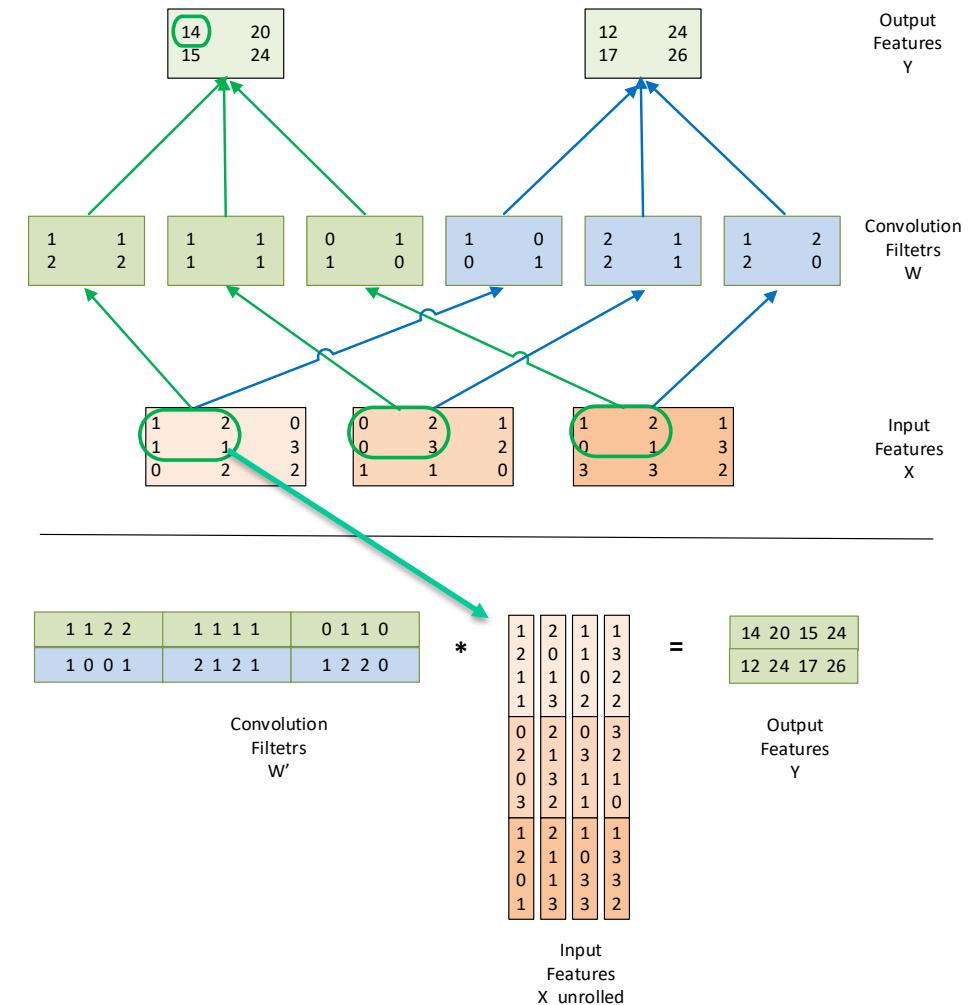
Output  
Features  
 $Y$

# Memory Access Efficiency of Original Convolution Algorithm

- Assume that we use tiled 2D convolution
- For input elements
  - Each output tile has  $\text{TILE\_WIDTH}^2$  elements
  - Each input tile has  $(\text{TILE\_WIDTH}+K-1)^2$
  - The total number of input feature map element accesses was  $\text{TILE\_WIDTH}^2 \times K^2$
  - The reduction factor of the tiled algorithm is  $K^2 \times \text{TILE\_WIDTH}^2 / ((\text{TILE\_WIDTH}+K-1)^2)$
- The convolution filter weight elements are reused within each output tile

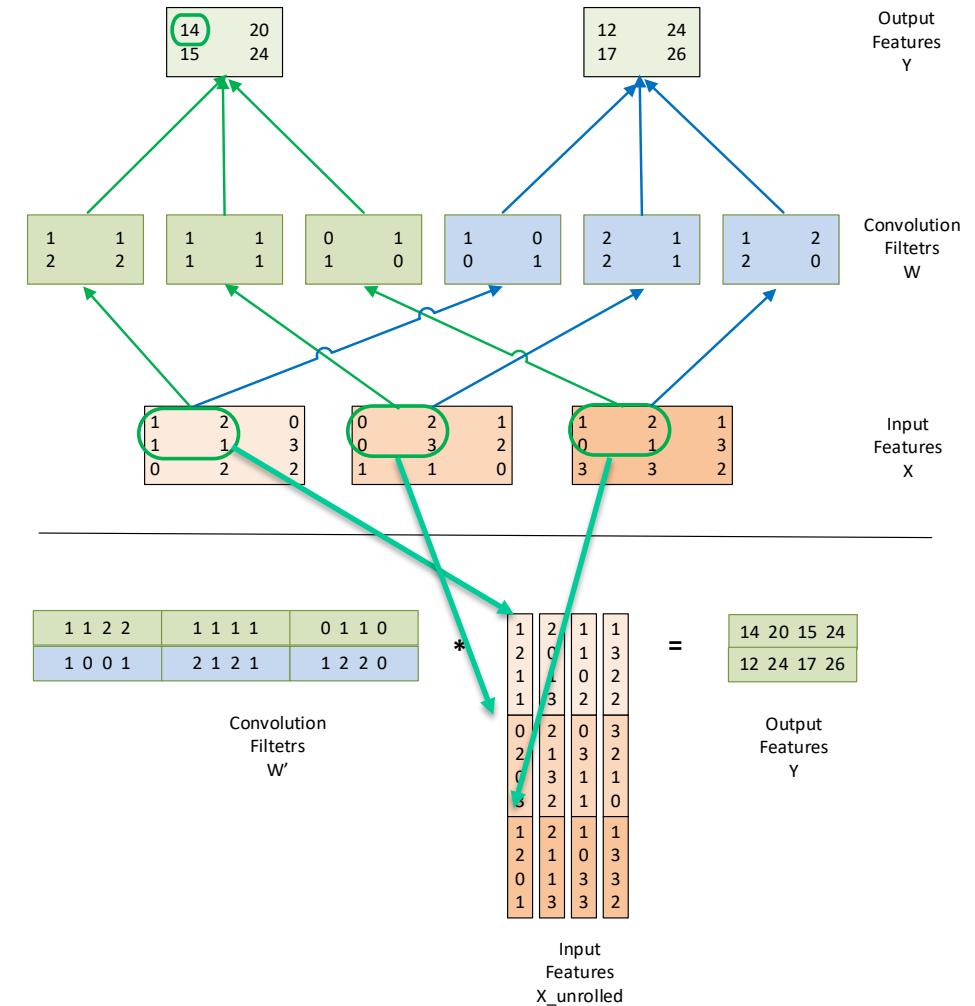
# Properties of the Unrolled Matrix

- Each unrolled column corresponds to an output feature map element
- For an output feature element  $(h, w)$ , the index for the unrolled column is  $h * W_{out} + w$  (linearized index of the output feature map element)



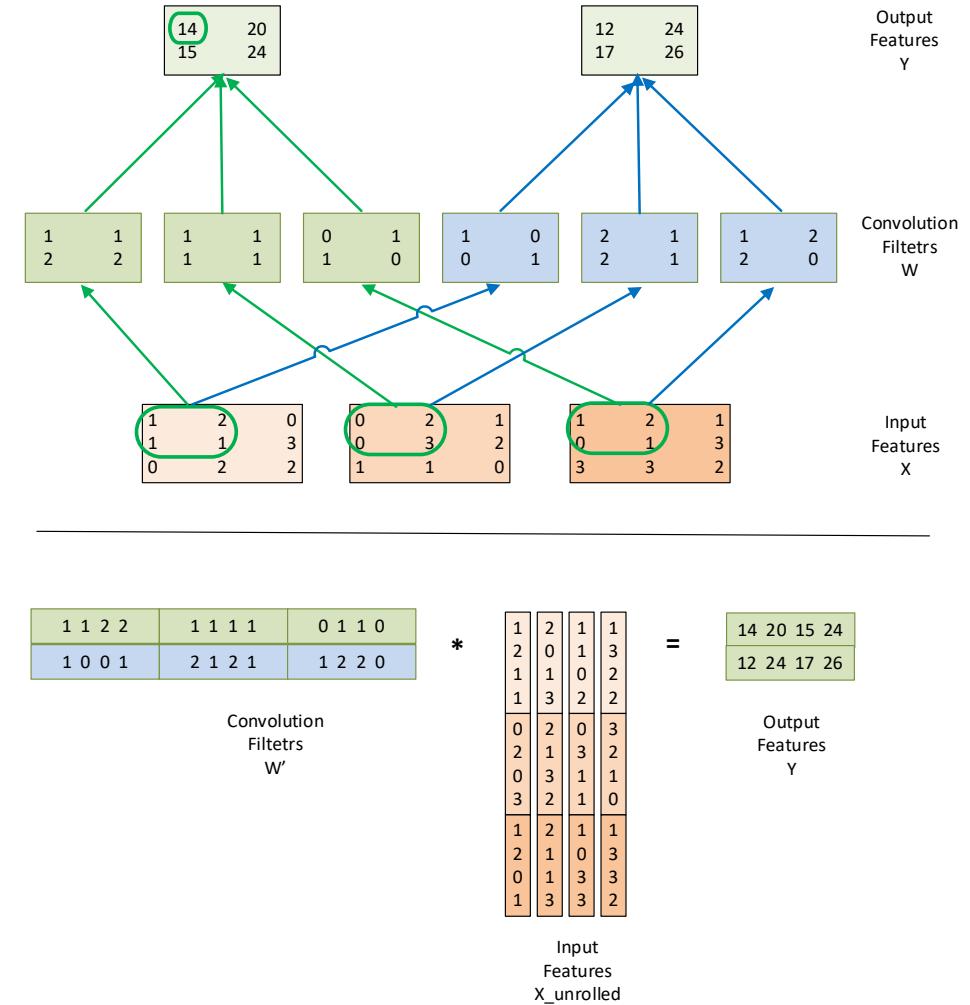
# Properties of the Unrolled Matrix (cont.)

- Each section of the unrolled column corresponds to an input feature map
- Each section of the unrolled column has  $k^*k$  elements (convolution mask size)
- For an input feature map  $c$ , the vertical index of its section in the unrolled column is  $c^*k^*k$  (linearized index of the output feature map element)



# To Find the Input Elements

- For output element  $(h, w)$ , the base index for the upper left corner of the input feature map  $c$  is  $(c, h, w)$
- The input element index for multiplication with the convolution mask element  $(p, q)$  is  $(c, h+p, w+q)$



# Input to Unrolled Matrix Mapping

Output element (h, w)

Mask element (p, q)

Input feature map c

// calculate the horizontal matrix index

```
int w_unroll = h * W_out + w;
```

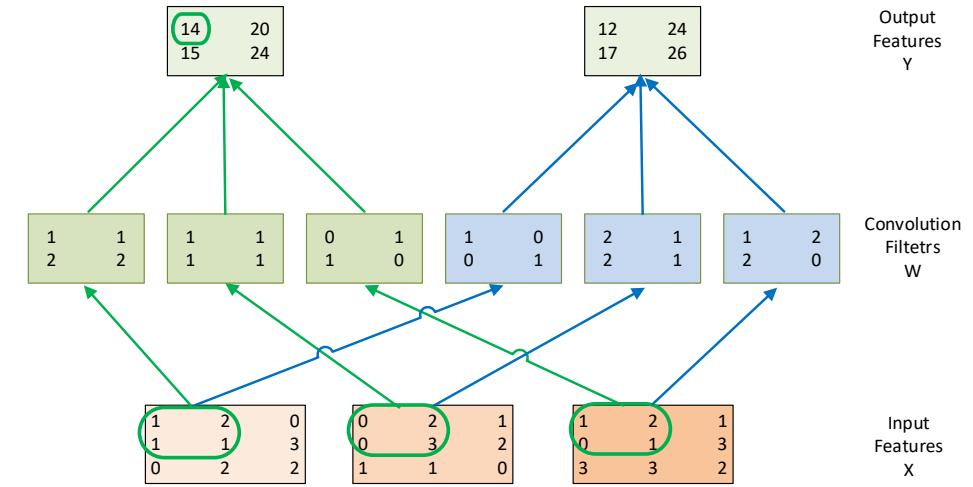
// find the beginning of the unrolled

```
int w_base = c * (K*K);
```

// calculate the vertical matrix index

```
int h_unroll = w_base + p * K + q;
```

```
X_unroll[b, h_unroll, w_unroll] = X[b, c, h + p, w + q];
```



$$\begin{array}{c} \begin{array}{|c|c|c|} \hline 1 & 1 & 2 & 2 \\ \hline 1 & 0 & 0 & 1 \\ \hline \end{array} & \begin{array}{|c|c|c|} \hline 1 & 1 & 1 & 1 \\ \hline 2 & 1 & 2 & 1 \\ \hline \end{array} & \begin{array}{|c|c|} \hline 0 & 1 & 1 & 0 \\ \hline 1 & 2 & 2 & 0 \\ \hline \end{array} \\ \times & \quad & \quad \\ \begin{array}{|c|c|c|} \hline 1 & 2 & 1 & 1 \\ \hline 2 & 0 & 1 & 3 \\ \hline 1 & 1 & 0 & 2 \\ \hline 1 & 3 & 2 & 2 \\ \hline 0 & 2 & 0 & 3 \\ \hline 2 & 1 & 3 & 2 \\ \hline 0 & 3 & 1 & 1 \\ \hline 3 & 2 & 1 & 0 \\ \hline 1 & 2 & 1 & 1 \\ \hline 2 & 1 & 0 & 3 \\ \hline 0 & 1 & 3 & 3 \\ \hline 1 & 3 & 3 & 2 \\ \hline \end{array} & = & \begin{array}{|c|c|c|} \hline 14 & 20 & 15 & 24 \\ \hline 12 & 24 & 17 & 26 \\ \hline \end{array} \\ \text{Convolution} \\ \text{Filters} \\ W' \\ \text{Input} \\ \text{Features} \\ X_{\text{unrolled}} \\ \end{array} \quad \begin{array}{c} \text{Output} \\ \text{Features} \\ Y \\ \end{array}$$

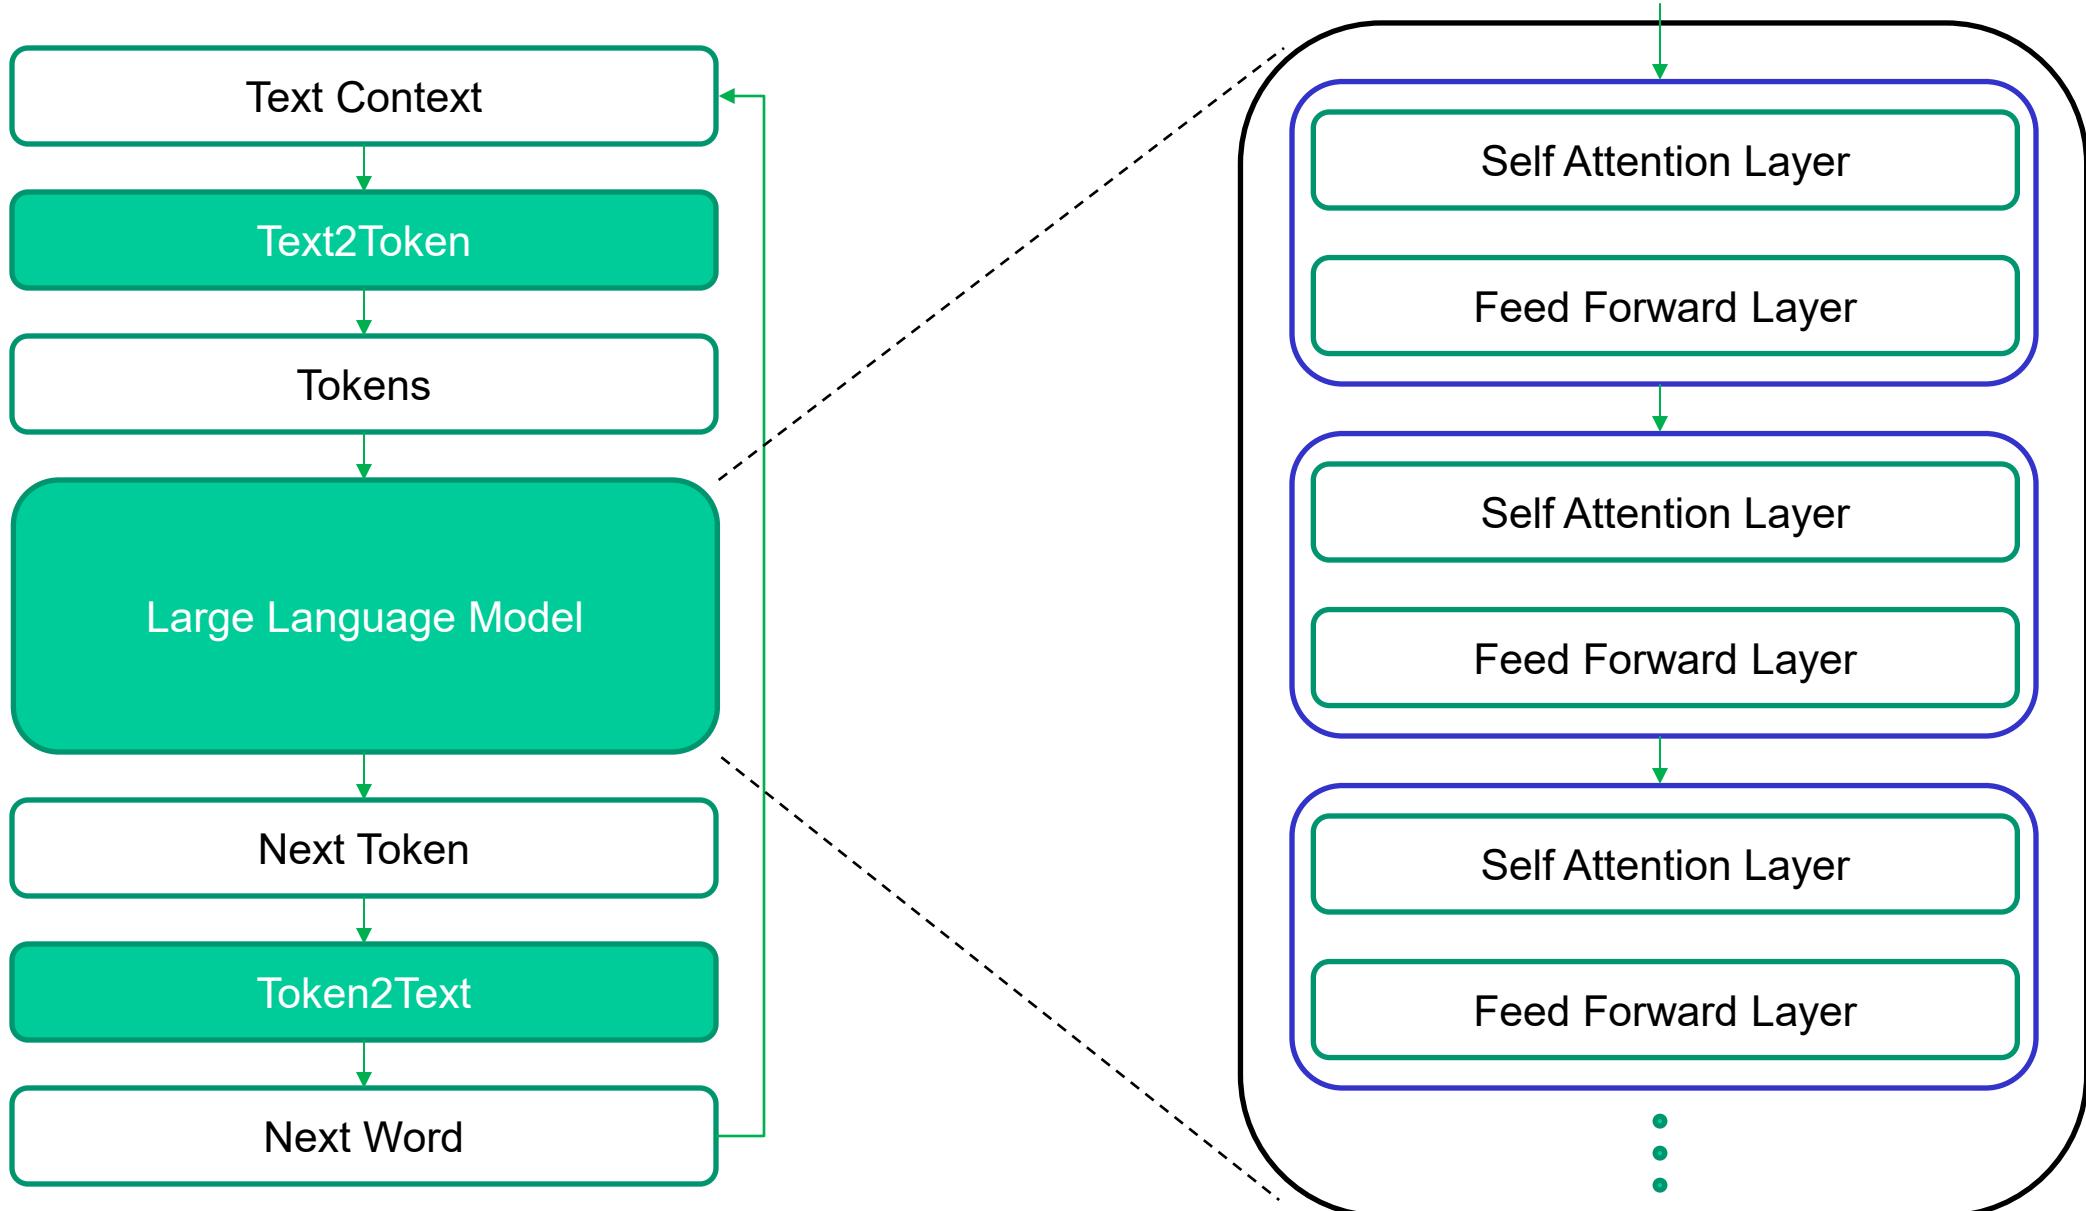
# Function to generate “unrolled” X

```
void unroll(int B, int C, int H, int W, int K, float* X, float* X_unroll)
{
    int H_out = H - K + 1;                                // calculate H_out, W_out
    int W_out = W - K + 1;
    for (int b = 0; b < B; ++b)                          // for each image
        for (int c = 0; c < C; ++c) {                     // for each input channel
            int w_base = c * (K*K);                      // per-channel offset for smallest X_unroll index
            for (int p = 0; p < K; ++p) {                  // for each element of KxK filter (two loops)
                for (int q = 0; q < K; ++q) {
                    for (int h = 0; h < H_out; ++h)        // for each thread (each output value, two loops)
                        for (int w = 0; w < W_out; ++w) {
                            int h_unroll = w_base + p * K + q; // data needed by one thread
                            int w_unroll = h * W_out + w;      // smallest index--across threads (output values)
                            X_unroll[b, h_unroll, w_unroll] = X[b, c, h + p, w + q]; // copy input pixels
                        }
                }
            }
        }
}
```

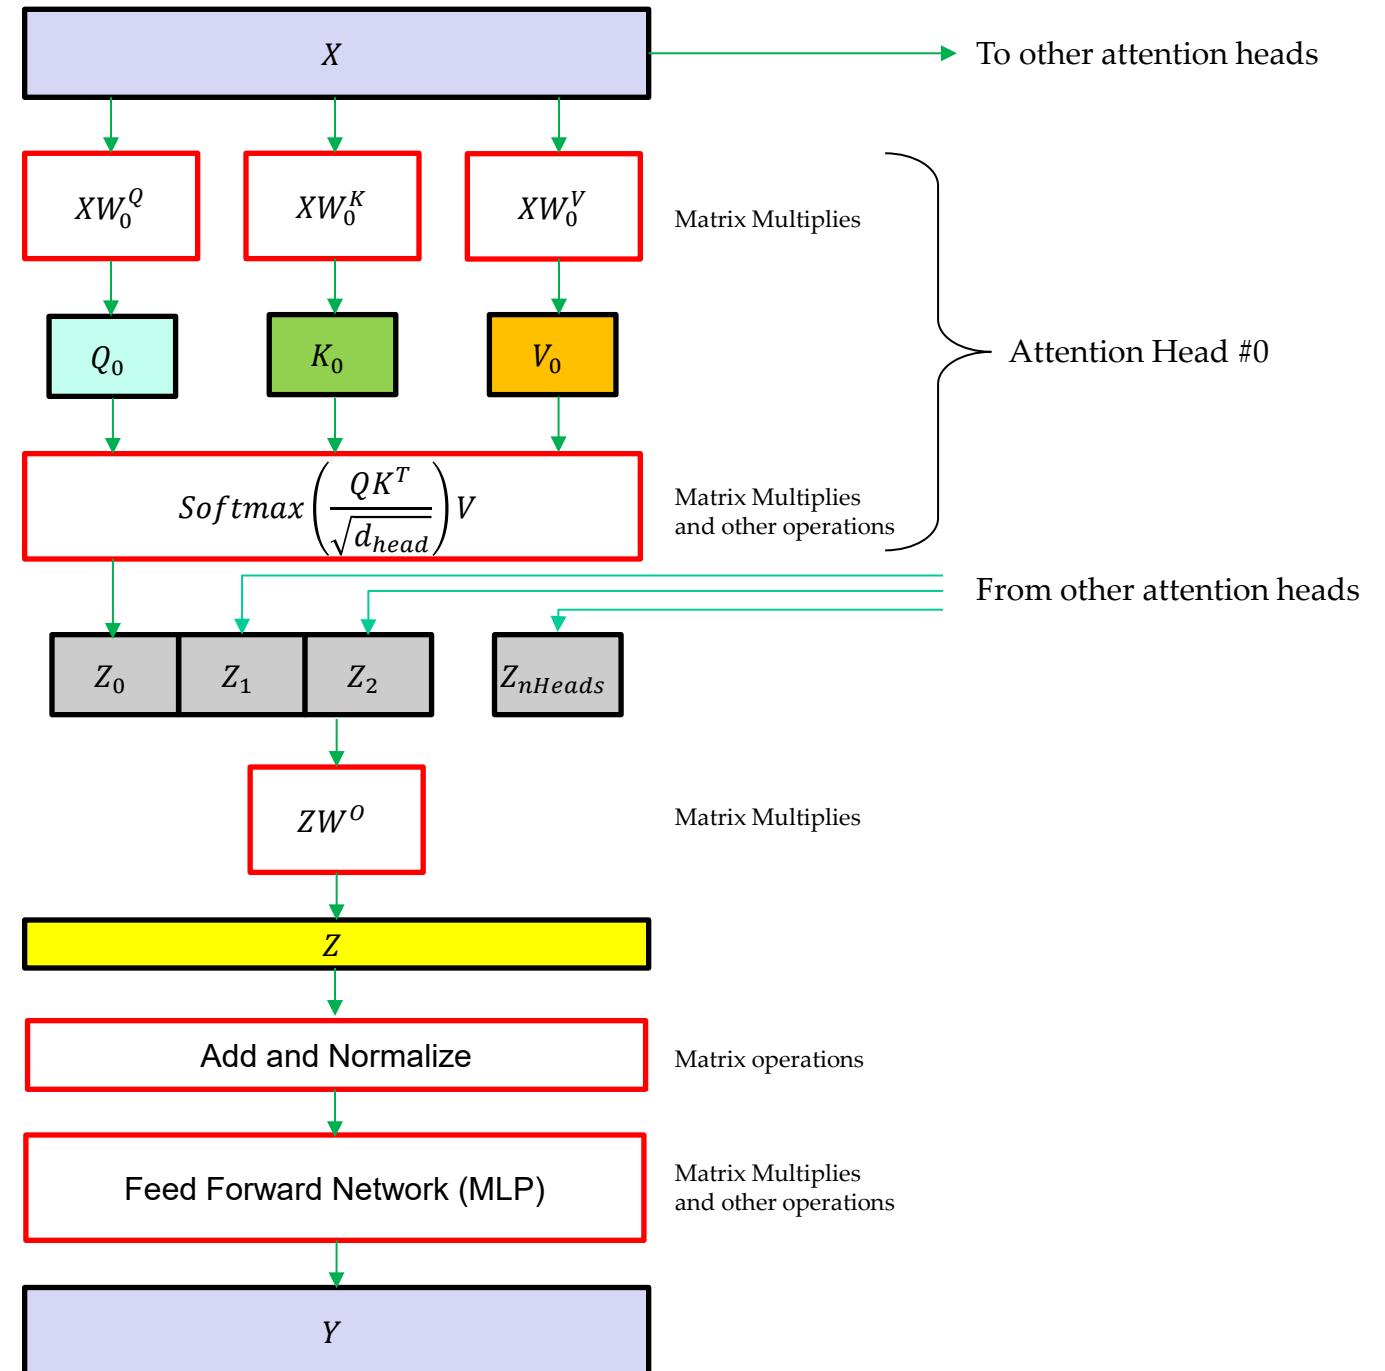
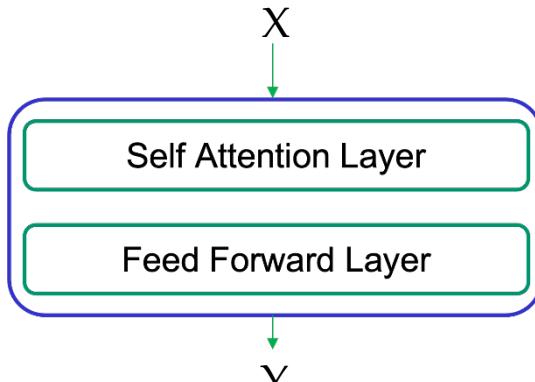
# Implementation Strategies for a Convolution Layer

- **Baseline**
  - Tiled 2D convolution implementation, use constant memory for convolution masks
- **Matrix-Multiplication Baseline**
  - Input feature map unrolling kernel, constant memory for convolution masks as an optimization
  - Tiled matrix multiplication kernel
- **Matrix-Multiplication with built-in unrolling**
  - Perform unrolling only when loading a tile for matrix multiplication
  - The unrolled matrix is only conceptual
  - When loading a tile element of the conceptual unrolled matrix into the shared memory, use the properties in the lecture to load from the input feature map
- **More advanced Matrix-Multiplication**
  - Use joint register-shared memory tiling

# Transformer-based Language Models



# Single Layer Computational Flow



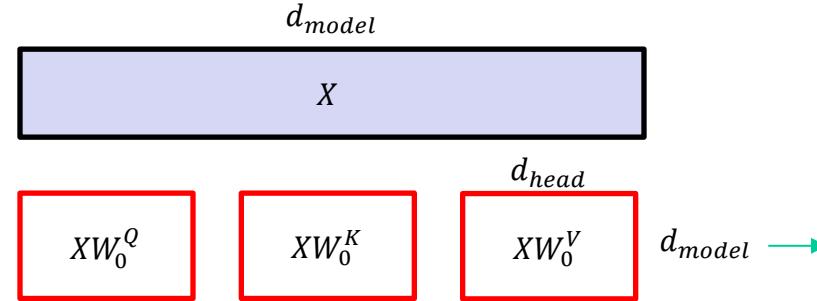
# GPT-3, as an example

Language Models are Few-Shot Learners, Brown et al., OpenAI, July 2020

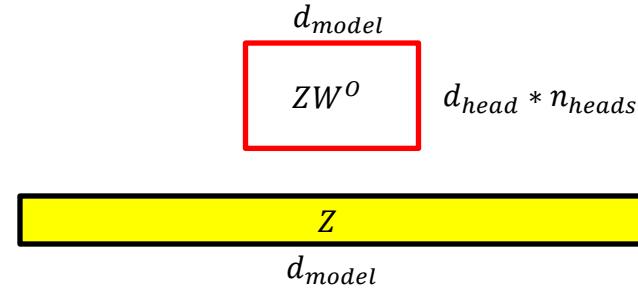
Model Name	$n_{\text{params}}$	$n_{\text{layers}}$	$d_{\text{model}}$	$n_{\text{heads}}$	$d_{\text{head}}$	Batch Size	Learning Rate
GPT-3 Small	125M	12	768	12	64	0.5M	$6.0 \times 10^{-4}$
GPT-3 Medium	350M	24	1024	16	64	0.5M	$3.0 \times 10^{-4}$
GPT-3 Large	760M	24	1536	16	96	0.5M	$2.5 \times 10^{-4}$
GPT-3 XL	1.3B	24	2048	24	128	1M	$2.0 \times 10^{-4}$
GPT-3 2.7B	2.7B	32	2560	32	80	1M	$1.6 \times 10^{-4}$
GPT-3 6.7B	6.7B	32	4096	32	128	2M	$1.2 \times 10^{-4}$
GPT-3 13B	13.0B	40	5140	40	128	2M	$1.0 \times 10^{-4}$
GPT-3 175B or "GPT-3"	175.0B	96	12288	96	128	3.2M	$0.6 \times 10^{-4}$

**Table 2.1:** Sizes, architectures, and learning hyper-parameters (batch size in tokens and learning rate) of the models which we trained. All models were trained for a total of 300 billion tokens.

GPT-3 has 96 Layers, 55B parameters just from Self Attention. 220 Gflop per X vector, per output token.



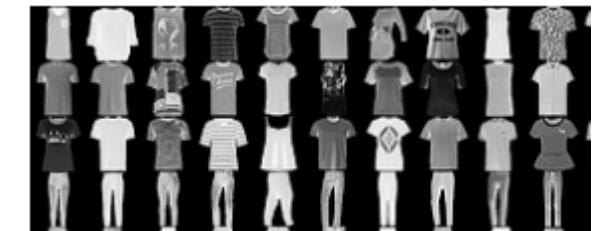
In GPT-3, each  $W$  is 12288x128 matrix, 1.5M parameters, 3MB (@ 16 bit floats), or 6 MFlop per  $X$  vector



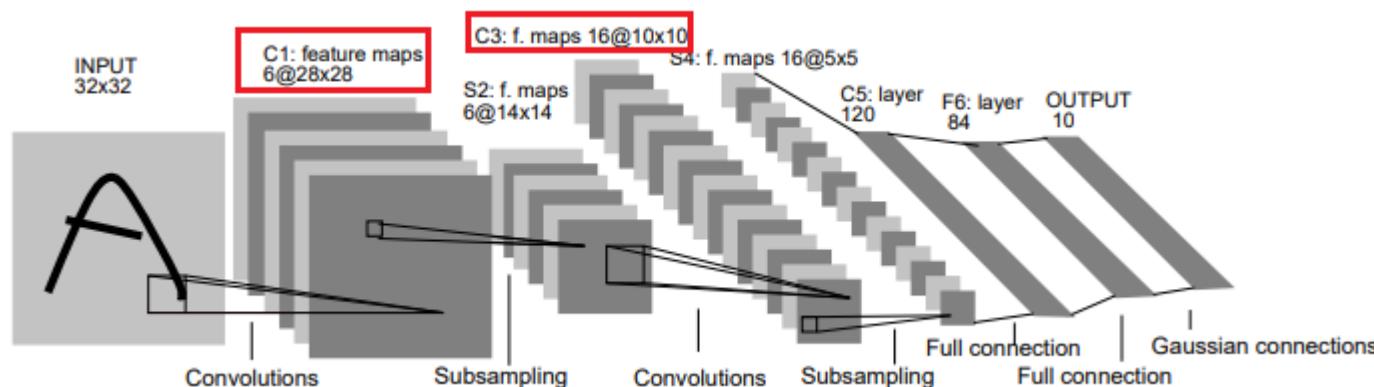
$W^O$  is 12288x12288 matrix, 150M parameters, 300MB (@ 16 bit floats), or 600 MFlop per  $X$  vector

# Project Overview

- Optimize the forward pass of the convolutional layers in a modified LeNet-5 CNN using CUDA. (CNN implemented using Mini-DNN, a C++ framework)
- The network will be classifying Fashion MNIST dataset
- Some network parameters to be aware of
  - Input Size: 86x86 pixels, batch of 1k-10k images
  - Input Channels: 1
  - Convolutional kernel size: 7x7
  - Number of kernels: Variable (your code should support this)



<https://github.com/zalandoresearch/fashion-mnist>



<http://yann.lecun.com/exdb/publis/pdf/lecun-01a.pdf>

# Project Timeline

- **All milestones are due on Fridays at 8 pm Central Time**
- Everyone must individually submit all milestones.
  - No sharing of code is allowed
- Project milestone 1:
  - CPU/GPU Convolution, code profiling
- Project milestone 2:
  - Unrolled matrix Convolution Kernel
- Project milestone 3:
  - GPU Convolution Kernel Optimizations

# Project Release

- Project is already released (only PM1 for now)
  - In the course github repository
  - [https://github.com/illinois-cs-coursework/f124\\_ece408\\_.release/tree/main/Project](https://github.com/illinois-cs-coursework/f124_ece408_.release/tree/main/Project)
- README.md in the repository contains all the instructions and details to complete the project.
- The github repo will be updated with additional code and instructions for PM2 & PM3
- Besides the code, there is a corresponding “report” to fill in on Canvas.



**ANY MORE QUESTIONS?  
READ CHAPTER 16**

# Problem Solving

- Q: Consider a 3D video filtering (convolution) code in CUDA with a  $5 \times 5 \times 7$  mask, which is stored in constant memory. Shared memory is used to fully store the input tile required for an  $16 \times 16 \times 16$  output tile (for example, using strategy 2). What is the ratio of total global memory read operations to shared memory accesses for one output tile? For this question, only consider interior tiles with no ghost elements.
- A:
  - $20 \times 20 \times 22$  to  $5 \times 5 \times 7 \times 16 \times 16 \times 16$

# Problem Solving

- Q: Consider a convolutional neural network that takes 100x200 images with 3 color channels (red, green, blue). The first layer of this network generates 10 output feature maps using 9x9 filters, where all channels are combined in each output feature map. Assuming all convolutions are performed in floating point, and considering only the convolutional layer (e.g., no pooling, thresholding, non-linearity, etc.), how many floating-point operations (both multiplications and additions) are required to generate all the output feature maps in a single forward pass? Remember: output feature maps are smaller than input maps because only pixels without ghost elements are generated.
- A:
  - 10 (output feature maps)
  - \* 92\*192 (total # of pixels to compute in each output feature map)
  - \* 9\*9 (conv. filter size)
  - \* 3 (color channels)
  - \* 2 (multiply-add operations)